

The background of the slide features a large, abstract cluster of colored dots in shades of orange, teal, pink, and yellow, arranged in a roughly circular pattern. A thin white curved line starts from the bottom left and sweeps upwards towards the top left corner.

# MongoDB: Coding for beginners

Macy Kung

# Welcome to the Course

- Introduction to NoSQL and MongoDB
- Knowledge Notes
- Set Up & Installation
- Inserting Documents
- Mongo Query Language
- Deleting Documents
- Updating Documents
- Data Modelling
- Index and Performance
- Aggregation

# Introduction to NoSQL Database

What is a Database?

What is an SQL Database?

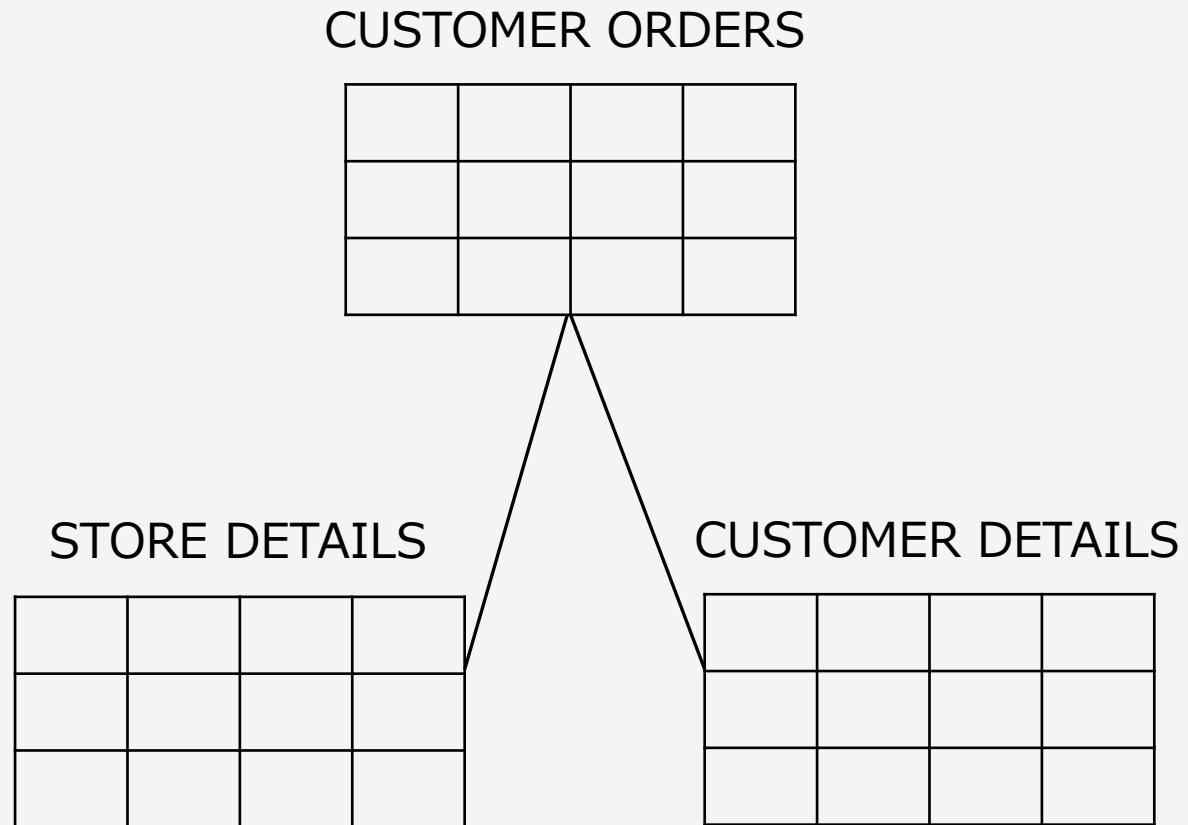
What is a NoSQL Database?



# **What is a Database**

A database is defined as an organized collection of structured information or data typically stored electronically.

# What is an SQL Database

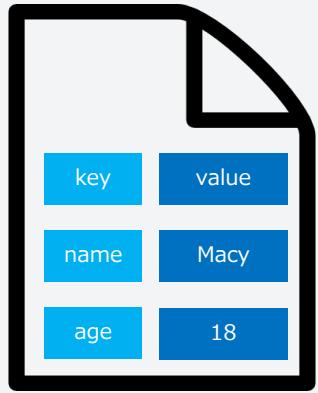


# What is a NoSQL Database

NoSQL (also refers to Not only SQL, non-SQL or non-relational) is a database which gives you a way to manage the data which is in a non-relational form i.e. which is not structured in a tabular manner and does not possess tabular relationships.

# Types of NoSQL Database

Document



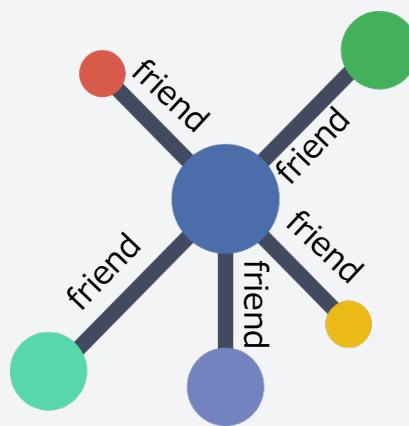
Apache, MongoDB, CouchDB

Key Value

key	value
Student_1	name@Laura @A4-tel@45789
Student_2	name@Chirs @B7-tel@12349

Redis, Dynamo, ZooKeeper

Graph



Neo4j, AllegroGraph,  
InfiniteGraph

Column Family

COLUMN FAMILY: Users			
Row Key	Columns		
Macy	name	age	email
			macy@...
James	name	age	email
			james@...

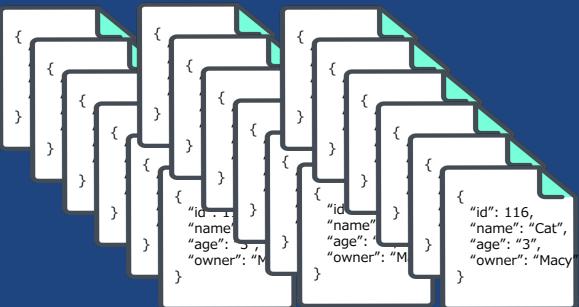
COLUMN FAMILY: Schools

Cassandra, Hbase, Google BigTable

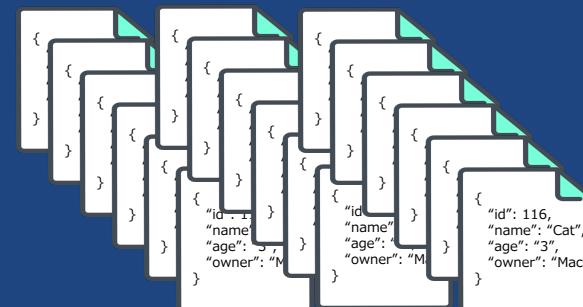
# What is MongoDB

DATABASE

COLLECTION

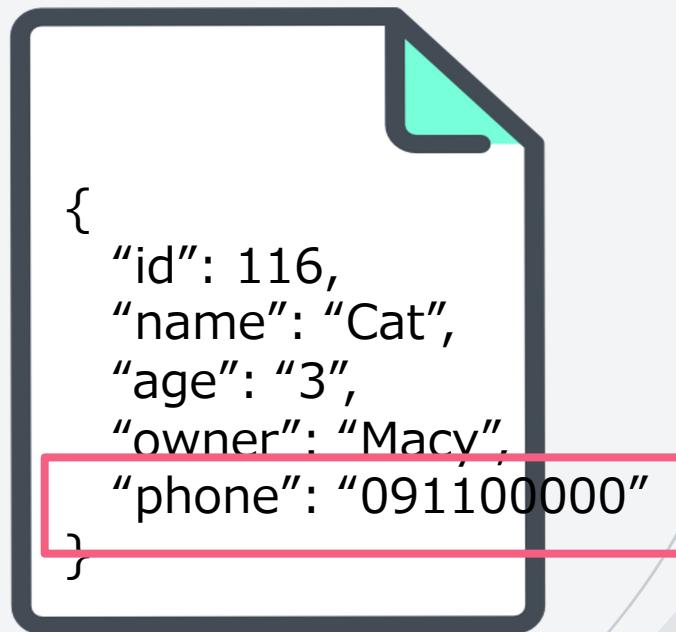
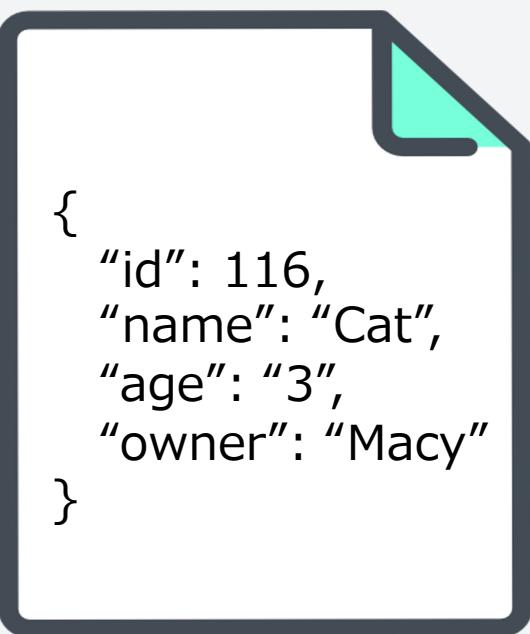


COLLECTION



```
{  
  "field": "value",  
  "field": "value",  
  "field": "value",  
  "field": "value"  
}
```

# What is MongoDB



Documents are **polymorphic** i.e. they don't have a fixed structure.  
Changes can be made easily to an individual document structure  
e.g. new field value pairs can be added.

# Previous example in SQL database

SQL

ID	NAME	AGE	OWNER
1	Cookie	3	Macy
2	Cin cin	7	Cindy
3	Dao	6	Lien
4	Pi	1	Yao



ID	NAME	AGE	OWNER	PHONE
1	Cookie	3	Macy	098100
2	Cin cin	7	Cindy	null
3	Dao	6	Lien	null
4	Pi	1	Yao	091888

# What is MongoDB

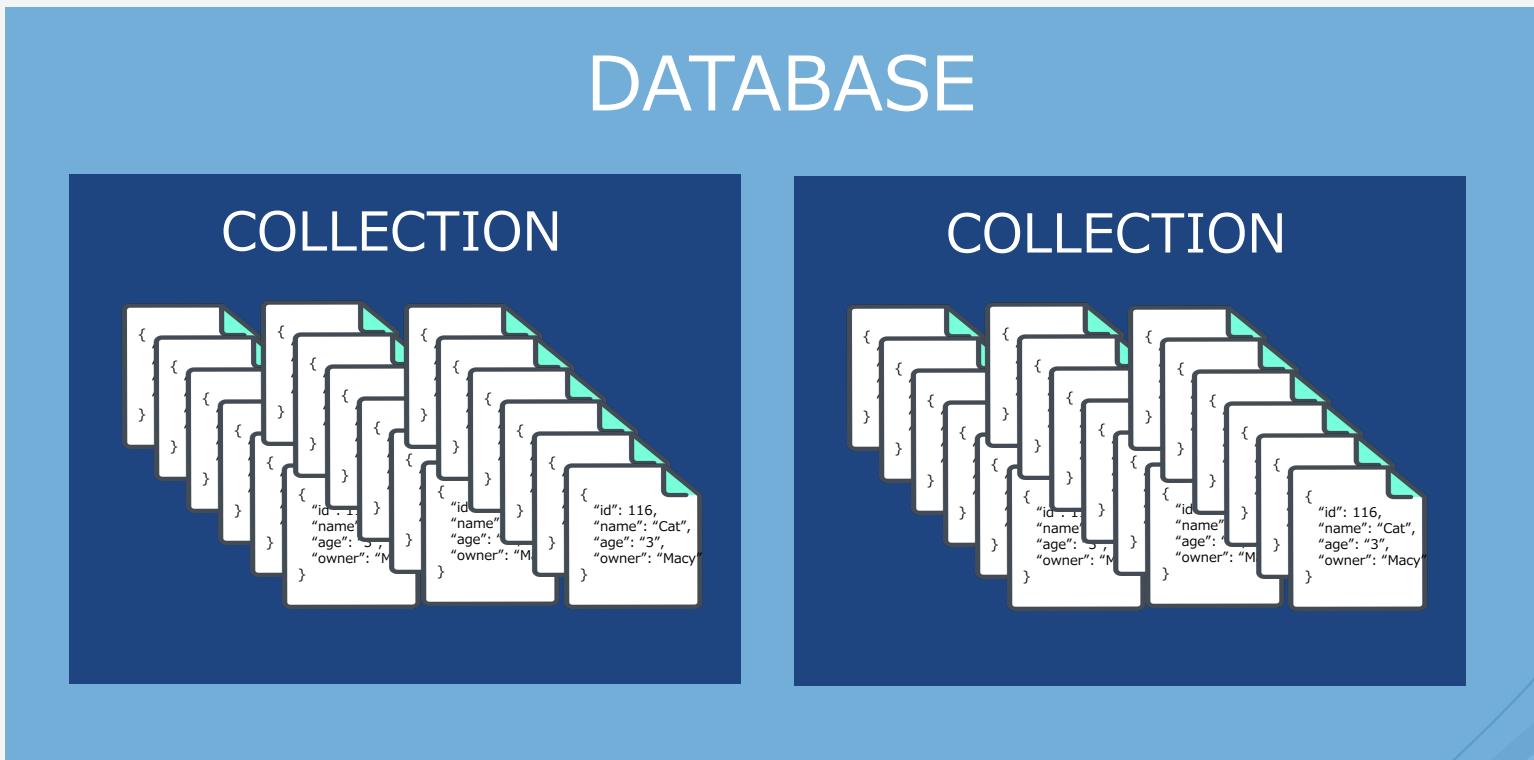
In general, data is stored more efficiently in MongoDB than in SQL because it is schema-less resulting in less data redundancy and more efficiency



# SQL vs NoSQL: comparison table



# Documents in MongoDB



# Documents in MongoDB



# Documents in MongoDB

```
{  
  "_id": ObjectId(11666),  
  "name": "Cat",  
  "address": {  
    "street": "No.1 Neihu road",  
    "city": "Taipei",  
    "postal code": "114"  
  }  
}
```

Nested Document

# Documents in MongoDB



```
{  
    "_id": ObjectId(11666),  
    "name": "Cat",  
    "phone": ["02-0034241", "091133442"]  
    "address": {  
        "street": "No.1 Neihu road",  
        "city": "Taipei",  
        "postal code": "114"  
    }  
}
```

An array is a data structure consisting of a collection of elements  
With an array data type, we can store multiple values in a single key of the document.

# Documents in MongoDB



```
{  
  "_id": ObjectId(11666),  
  "name": "Cat",  
  "age": "3",  
  "owner": "Macy"  
}
```



```
{  
  "_id": ObjectId(11666),  
  "name": "Cat",  
  "age": "3",  
  "owner": "Macy",  
  "phone": "0911000000"  
}
```



```
{  
  "_id": ObjectId(11666),  
  "name": "Cat",  
  "age": "3",  
  "owner": "Macy",  
  "cell phone": "0911000000"  
}
```

Not the same field

```
{"hello": "world"} → \x16\x00\x00\x00          // total document size  
 \x02                      // 0x02 = type String  
 hello\x00                  // field name  
 \x06\x00\x00\x00world\x00    // field value  
 \x00                      // 0x00 = type E00 ('end of object')
```

```
{"BSON": ["awesome", 5.05, 1986]} → \x31\x00\x00\x00  
 \x04BSON\x00  
 \x26\x00\x00\x00  
 \x02\x30\x00\x08\x00\x00\x00awesome\x00  
 \x01\x31\x00\x33\x33\x33\x33\x33\x33\x33\x14\x40  
 \x10\x32\x00\xc2\x07\x00\x00  
 \x00  
 \x00
```

# What is BSON? Does MongoDB use BSON, or JSON?

# MongoDB - Datatypes

**NUMERICAL:**  
Numbers such as integers  
and decimals

**DATE:**  
Date and time values

**STRING:**  
Charater / text  
values

👉 [Datatypes Reference](#)

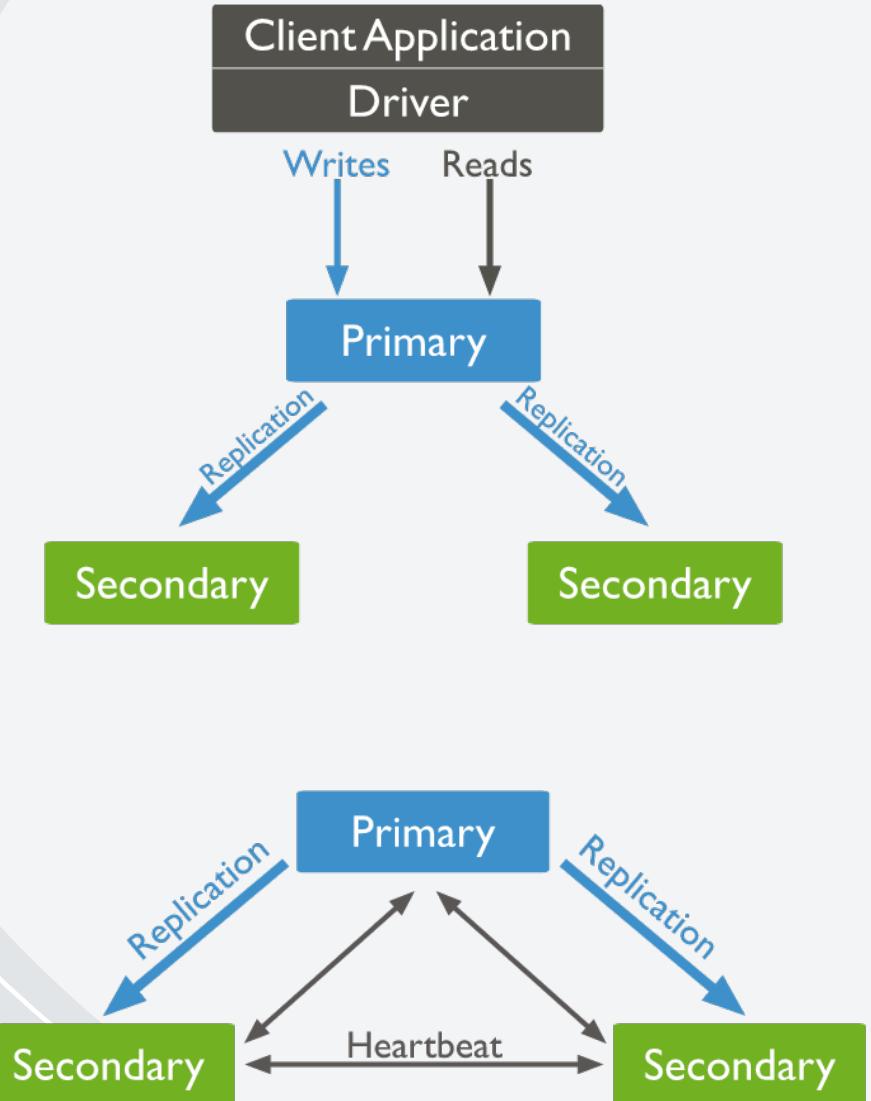
# Replication

- A replica set in MongoDB is a group of **mongod** processes that maintain the same data set. Replica sets provide redundancy and high availability, and are the basis for all production deployments.
- There are 3 instances in a replica set and each instance contains a full copy of the data in each instance



# Replication

- Of the data bearing nodes, one and only one member is deemed the primary node, while the other nodes are deemed secondary nodes. If the primary is unavailable, an eligible secondary will hold an election to elect itself the new primary.
- **mongod** is the primary daemon process for the MongoDB system. It handles data requests, manages data access, and performs background management operations.



# Sharded Cluster

- Sharding is a method for distributing data across multiple machines. MongoDB uses sharding to support deployments with very large data sets and high throughput operations.



# Sharded cluster consists of the following components

## SHARD

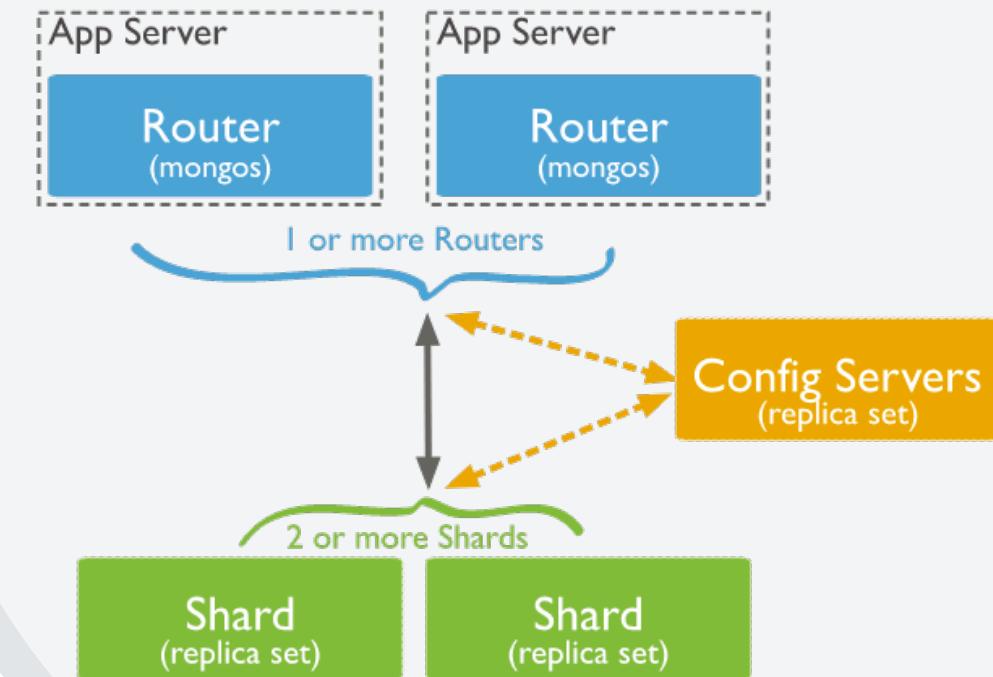
Containing a partition of the data

## MONGOS

Interface between any application using MongoDB and the sharded cluster

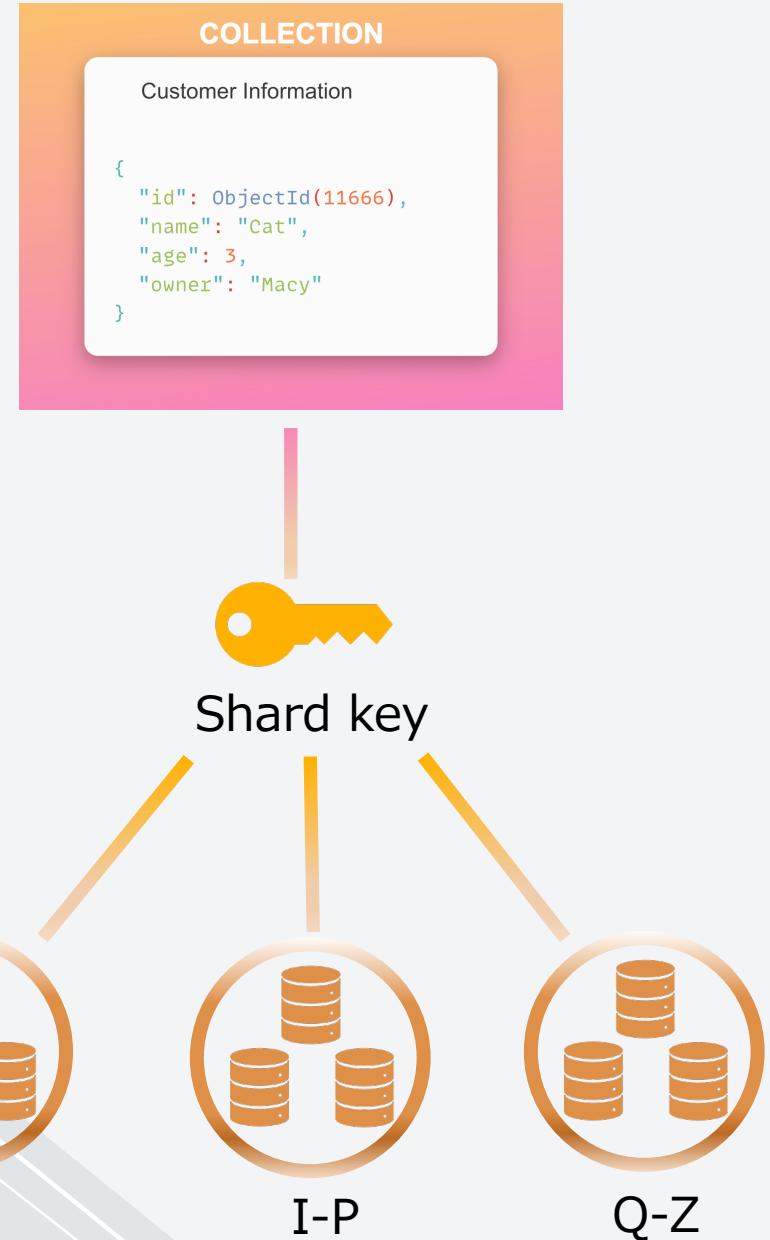
## CONFIG SERERS

Store metadata and configuration settings



# Natively Distributed System

- MongoDB uses the **shard key** to distribute the collection's documents across shards.
- The shard key consists of a field or multiple fields in the documents.



# You get a basic knowledge of MongoDB

- What's NoSQL database
- MongoDB with format of document
- A replica set typically consists of \_\_ machine(s)?
- A typical Relational Database allows you more flexibility in data structure than MongoDB

# Advantages of MongoDB Database

- Less complexity vs relational databases
- Schema-less
- Easier to maintain
- Easy to scale

# When to use NoSQL

- To handle a huge volume of structured, semi-structured and unstructured data.
- If your relational database is not capable enough to scale up to your traffic at an acceptable cost.
- If you want to have an efficient, scale-out architecture in place of an expensive and monolithic architecture.
- If you have local data transactions that need not be very durable.
- If you are going with schema-less data and want to include new fields without any ceremony.
- When your priority is easy scalability and availability.

# When to Avoid NoSQL

- In such cases like financial transactions, etc., you may go with SQL databases.
- If consistency is a must and if there aren't going to be any large-scale changes in terms of the data volume, then going with the SQL database is a better option.
- One should also keep in mind that NoSQL databases won't support structured query language. The querying language may vary from one database to another.

# Set Up & Installation

- Click Try Free, SignUp to [MongoDB Atlas](#) with Google Account
- Download [Mongo Compass](#)
- MongoDB Atlas Quick Tour
- Create database
- Create Users
- Network access
- Connect to MongoDB using Mongo Compass
- [Import data](#)

# Inserting Documents - insertOne



insertOne

```
db.collection.insertOne(  
  <document or array of documents>  
)
```

Inserts a single document into a collection.  
Returns an object that contains  
the status of the operation.

Collection: shoes

```
db.shoes.insertOne({  
  brand: "Nike",  
  item: "Fragment x Sacai x Nike LDWaffle",  
  price: 15000,  
  color: "gray",  
  size: "UK5"  
})
```

Return :

```
{  
  "_id": ObjectId("537a41a"),  
  "brand": "Nike",  
  "item": "Fragment x Sacai x Nike LDWaffle",  
  "price": 15000,  
  "color": "gray",  
  "size": "UK5"  
}
```

# Inserting Documents - insertMany



insertMany

```
db.collection.insertMany(  
  <document or array of documents>  
)
```

Inserts multiple documents  
into a collection.



Collection: shoes

```
db.shoes.insertMany([  
  { brand: "Nike",  
    item: "Fragment X Sacai X Nike LDWaffle",  
    price: 13000,  
    color: "blue",  
    size: "UK5"  
  },  
  { brand: "Nike",  
    item: "Fragment X Sacai X Nike LDWaffle",  
    price: 15000,  
    color: "blue",  
    size: "UK8"  
  },  
])
```

# Using the find Methods



find

```
db.collection.find({query},{projection})
```

Returns all documents in a collection

Query and projection are optional

- Query : Specifies selection filter.
- Projection: Specifies the fields to return in the documents that match the query filter.
  - 1: true, query this field
  - 0: false, ignore this field

Collection: shoes

```
[  
  {_id:ObjectId("11ssa2")}, brand: "Hoka", size: "UK7"},  
  {_id:ObjectId("0a4sde")}, brand: "Nike", size: "UK7"},  
  {_id:ObjectId("ya85de")}, brand: "Hoka", size: "UK5"}  
]
```

find

MongoQL: db.shoes.find({"brand":"Hoka"}, {"size":1})

SQL : select \_id, size from shoes where brand='Hoka'

Return :

```
[  
  {_id:ObjectId("11ssa2")}, size: "UK7"},  
  {_id:ObjectId("ya85de")}, size: "UK5"}  
]
```

# Using the findOne Methods



findOne

```
db.collection.findOne({query},{projection})
```

Returns one document from the collection  
If multiple documents satisfy the query,  
the method returns the first document.

Collection: shoes

```
[  
  {_id:ObjectId("11ssa2")}, brand: "Hoka", size: "UK7"},  
  {_id:ObjectId("0a4sde")}, brand: "Nike", size: "UK7"},  
  {_id:ObjectId("ya85de")}, brand: "Hoka", size: "UK5"}  
]
```

findOne

MongoQL: db.shoes.findOne({"brand": "Hoka"})

SQL : select \* from shoes where brand='Hoka' limit 1

Return:

```
{_id:ObjectId("11ssa2")}, brand: "Hoka", size: "UK7"}
```

# Case Sensitivity in MongoDB

Objects in MongoDB are case sensitive.

For example, the `findOne` method (performed in the `sample_training` database on the `shoes` collection) should be as follows:

**`db.shoes.findOne()`** – this is the correct syntax

**`Db.shoes.findOne()`** – will not work because `Db` contains an upper case character

**`db.Shoes.findOne()`** – will not work because the collection is saved as “`trips`” not “`Trips`”, case sensitivity matters

**`db.shoes.FindOne()`** – will not work because case sensitivity also matters on methods. It should be “`findOne`” not “`FindOne`”

The findOne method performed in the database on the shoes collection

- a. db.shoes.findOne()
- b. Db.shoes.findOne()
- c. db.Shoes.findOne()
- d. db.shoes.FindOne()
- e. All of the above



**Which of the  
following is correct**

# Use of \$ in Mongo query language

- Precedes operators in Mongo query language
- Precedes field values
- Aggregation Pipeline



Equal to

**{"size":{\$eq:"UK8"}}**

**{field:{operator : value}}**

# Logical Operators

\$and

Return all documents  
that match all conditions

\$or

Return all documents that  
match any of the conditions

\$nor

Return all documents that  
fail all of the conditions

{operator: [{condition1},{condition2},..]}

\$not

Return all documents that  
do not match the expression

{operator: [{condition1}]}  
}

# Comparison Operators

\$in

values specified in an array  
{"size":{\$in:["UK8", "UK11"]}}

\$lt

Less than

\$gt

Greater than

\$eq

Equal to  
{"size":"UK8"}  
{"size":{\$eq:"UK8"}}}

\$nin

Not in

\$gte

Greater than  
Or Equal to

\$lte

Less than  
Or equal to

\$ne

Not equal to

# Elements Operators

\$exists

Return documents that contain the specified field

**db.collection.find({field:{\$exists:<boolean>}})**

\$type

Return documents that contain values of a specified data type

**db.collection.find({field:{\$type:<BSON type>}})**

**db.collection.find({field:{\$type:[<BSON type1>, <BSON type2>,...]}})**

# Query Embedded Documents

You can perform read operations  
using the db.collection.find()  
method

- Match an Embedded/Nested Document
- Specify Equality Match on a Nested Field
- Specify Match using Query Operator
- Specify AND Condition

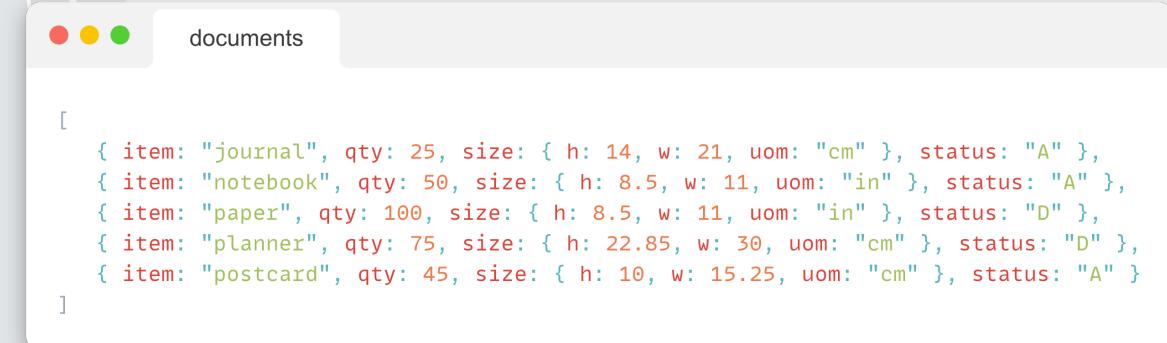
# Match an Embedded/Nested Document

To specify an equality condition on a field that is an embedded/nested document, use the [query filter document](#)

**{ <field>: <value>}**

where <value> is the document to match.

```
db.inventory.find({ size: { h: 14, w: 21, uom: "cm" } })
```

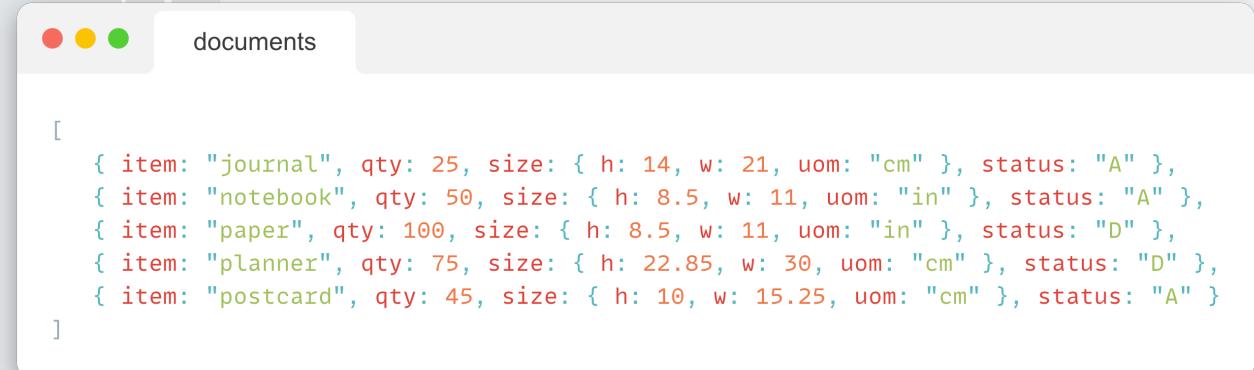


```
[{"item": "journal", "qty": 25, "size": { h: 14, w: 21, uom: "cm" }, "status": "A"}, {"item": "notebook", "qty": 50, "size": { h: 8.5, w: 11, uom: "in" }, "status": "A"}, {"item": "paper", "qty": 100, "size": { h: 8.5, w: 11, uom: "in" }, "status": "D"}, {"item": "planner", "qty": 75, "size": { h: 22.85, w: 30, uom: "cm" }, "status": "D"}, {"item": "postcard", "qty": 45, "size": { h: 10, w: 15.25, uom: "cm" }, "status": "A"}]
```

# Specify Equality Match on a Nested Field

The following example selects all documents where the field uom nested in the size field equals "in"

```
db.inventory.find({ "size.uom": "in" })
```



A screenshot of a MongoDB shell window titled "documents". The window shows a cursor of documents, indicated by a blue border around the list. The documents are represented as JSON objects:

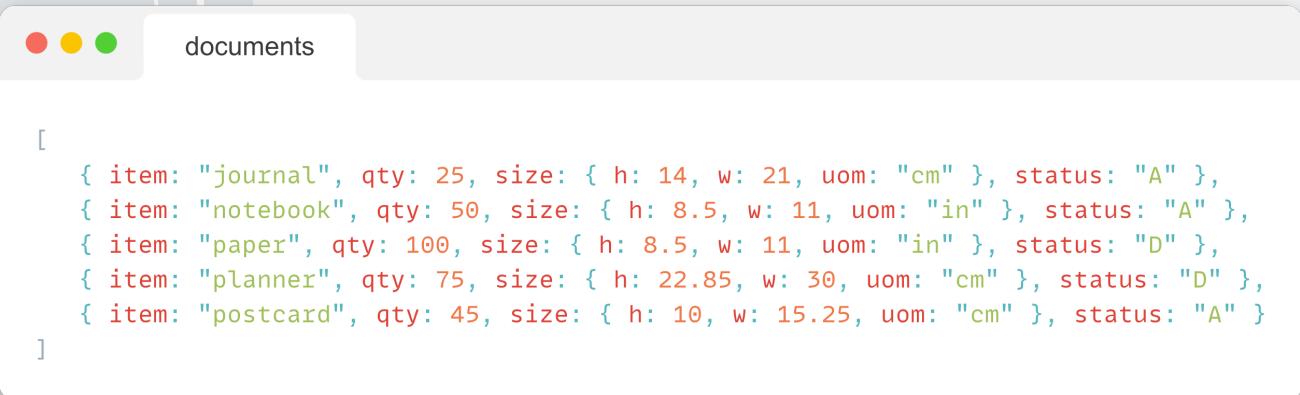
```
[  
  { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },  
  { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },  
  { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },  
  { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },  
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" }]
```

# Specify Match using Query Operator

A [query filter document](#) can use the [query operators](#) to specify conditions in the following form:

```
{ <field1>: { <operator1>:  
<value1> }, ... }
```

```
db.inventory.find({ "size.h": { $lt: 15 } })
```



The screenshot shows a MongoDB shell window with a title bar labeled "documents". The content area displays a JSON array representing the results of a query. The array contains five documents, each representing a different item with its quantity, size (height and width), uom (unit of measurement), and status.

```
[  
  { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },  
  { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },  
  { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },  
  { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },  
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" }]
```

# Specify AND Condition

The following query selects all documents where the nested field h is less than 15, the nested field uom equals "in", and the status field equals "D"

```
db.inventory.find({ "size.h": { $lt: 15 }, "size.uom": "in", status: "D" })
```

documents

```
[  
  { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },  
  { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },  
  { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },  
  { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },  
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" }  
]
```

# Query an Array of Embedded Documents

- Query for a Document Nested in an Array
- Specify a Query Condition on a Field Embedded in an Array of Documents
- Use the Array Index to Query for a Field in the Embedded Document
- A Single Nested Document Meets Multiple Query Conditions on Nested Fields
- Combination of Elements Satisfies the Criteria

# Query for a Document Nested in an Array

The following example selects all documents where an element in the `instock` array matches the specified document

```
db.inventory.find( { "instock": { warehouse: "A", qty: 5 } } )
```



A screenshot of a MongoDB shell window titled "documents". The window shows a list of documents in JSON format. The first document is highlighted with a red border. The query used to find this document is shown above the results.

```
[  
  { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] },  
  { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] },  
  { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] },  
  { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
)
```

# Specify a Query Condition on a Field Embedded in an Array of Documents

If you do not know the index position of the document nested in the array, concatenate the name of the array field, with a dot (.) and the name of the field in the nested document.

The following example selects all documents where the `instock` array has at least one embedded document that contains the field `qty` whose value is less than or equal to 20

```
db.inventory.find( { 'instock.qty': { $lte: 20 } } )
```

A screenshot of a MongoDB query results window. At the top, there's a title bar with three colored dots (red, yellow, green) and the word "documents". Below the title bar, the word "[ ]" indicates an empty array. Then, a list of five documents is shown, each representing an item in the inventory:

```
[  
  { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] },  
  { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] },  
  { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] },  
  { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
]
```

# Use the Array Index to Query for a Field in the Embedded Document

Using [dot notation](#), you can specify query conditions for field in a document at a particular index or position of the array. The array uses zero-based indexing.

The following example selects all documents where the `instock` array has as its first element a document that contains the field `qty` whose value is less than or equal to 20

```
db.inventory.find( { 'instock.0.qty': { $lte: 20 } } )
```



A screenshot of a MongoDB shell window titled "documents". The window shows the results of a query. The results are displayed in a JSON array, starting with a bracket "[", followed by five document objects. Each document has an "item" field and an "instock" array. The "instock" array contains two elements: a warehouse with a quantity of 5 and another warehouse with a quantity of 15. The "item" fields are "journal", "notebook", "paper", "planner", and "postcard". The "instock" arrays for these items are: [{"warehouse": "A", "qty": 5}, {"warehouse": "C", "qty": 15}], [{"warehouse": "C", "qty": 5}], [{"warehouse": "A", "qty": 60}, {"warehouse": "B", "qty": 15}], [{"warehouse": "A", "qty": 40}, {"warehouse": "B", "qty": 5}], [{"warehouse": "B", "qty": 15}, {"warehouse": "C", "qty": 35}]. The window has a standard OS X title bar with red, yellow, and green buttons.

```
[  
  { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] },  
  { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] },  
  { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] },  
  { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
)
```

# A Single Nested Document Meets Multiple Query Conditions on Nested Fields

Use \$elemMatch operator to specify multiple criteria on an array of embedded documents such that at least one embedded document satisfies all the specified criteria.

The following example queries for documents where the instock array has at least one embedded document that contains both the field qty equal to 5 and the field warehouse equal to A

```
db.inventory.find( { "instock": { $elemMatch: { qty: 5, warehouse: "A" } } } )
```

documents

```
[  
  { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] },  
  { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] },  
  { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] },  
  { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
)
```

The following example queries for documents where the instock array has at least one embedded document that contains the field qty that is greater than 10 and less than or equal to 20

```
db.inventory.find( { "instock": { $elemMatch: { qty: { $gt: 10, $lte: 20 } } } } )
```

documents

```
[  
  { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] },  
  { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] },  
  { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] },  
  { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
)
```

# Combination of Elements Satisfies the Criteria

If the compound query conditions on an array field do not use the [\\$elemMatch](#) operator, the query selects those documents whose array contains any combination of elements that satisfies the conditions.

For example, the following query matches documents where any document nested in the instock array has the qty field greater than 10 and any document (but not necessarily the same embedded document) in the array has the qty field less than or equal to 20

```
db.inventory.find( { "instock.qty": { $gt: 10, $lte: 20 } } )
```



The screenshot shows a macOS-style application window titled "documents". Inside, the MongoDB shell displays the results of a query. The results are an array of documents, each representing an item with its instock array. The array contains objects for different warehouses, some with qty values greater than 10 and others less than or equal to 20, thus satisfying the compound query condition.

```
[ { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] }, { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] }, { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] }, { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] }, { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }
```

The following example queries for documents where the instock array has at least one embedded document that contains the field qty equal to 5 and at least one embedded document (but not necessarily the same embedded document) that contains the field warehouse equal to A

```
db.inventory.find( { "instock.qty": 5, "instock.warehouse": "A" } )
```



A screenshot of a Mac OS X application window titled "documents". The window contains a list of inventory items, each with an "item" name and an "instock" array. The "instock" array contains multiple objects, each with a "warehouse" field and a "qty" field. The query filters for items where at least one "instock" object has a "qty" of 5 and at least one "instock" object has a "warehouse" of "A".

```
[  
  { item: "journal", instock: [ { warehouse: "A", qty: 5 }, { warehouse: "C", qty: 15 } ] },  
  { item: "notebook", instock: [ { warehouse: "C", qty: 5 } ] },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 }, { warehouse: "B", qty: 15 } ] },  
  { item: "planner", instock: [ { warehouse: "A", qty: 40 }, { warehouse: "B", qty: 5 } ] },  
  { item: "postcard", instock: [ { warehouse: "B", qty: 15 }, { warehouse: "C", qty: 35 } ] }  
)
```

# Deleting Documents - deleteOne



deleteOne

```
db.collection.deleteOne(<filter>)
```

The first document that matches  
the filter condition will be deleted

Collection: shoes

```
{  
  _id: ObjectId("537a41a") ,  
  brand: "Nike" ,  
  item: "Fragment x Sacai x Nike LDWaffle" ,  
  price: 15000 ,  
  color: "gray" ,  
  size: "UK5"  
}
```

deleteOne

MongoQL:

```
db.shoes.deleteOne({"_id": ObjectId("537a41a")})
```

Return :

```
{ "acknowledged" : true , "deletedCount" : 1 }
```

# Delete Documents - deleteMany



deleteMany

```
db.collection.deleteMany(<filter>)
```

All documents that match the condition will be deleted.

Collection: shoes

```
[  
  {  
    _id: ObjectId("537a41a"),  
    brand: "Nike",  
    item: "Fragment x Sacai x Nike LDWaffle",  
    price: 15000,  
    color: "grey",  
    size: "UK5"  
  }, ...]
```

deleteMany

MongoQL:

```
db.shoes.deleteMany({"color": "grey"})
```

Return :

```
{ "acknowledged" : true, "deletedCount" : 8 }
```

# Updating Documents



updateOne

```
db.collection.updateOne(  
  <filter>, <update>, <options>  
)
```

Update one will update the first documents that match the filter criteria.



updateMany

```
db.collection.updateMany(  
  <filter>, <update>, <options>  
)
```

Update many will update all documents that match the filter criteria.

\$set

**db.collection.updateOne({filter}, {  
 \$set:{field:value, ...},  
 options})**

**db.collection.updateMany({filter}, {  
 \$set:{field:value, ...},  
 options})**

# Update Operators

\$unset

Removes the specified field from a document.

**db.collection.updateMany({filter},{\$unset:{field:value,...}},options)**

\$inc

Increments the value of the field by the specified amount.

**db.collection.updateMany({filter},{\$inc:{field:value,...}},options)**

\$rename

Renames a field.

**db.collection.updateMany({filter},{\$rename:{field:value,...}},options)**

\$push

Adds an item to an array.

**db.collection.updateMany({filter},{\$push:{field:value,...}},options)**

Please refer to this link for additional information on [MongoDB Update Operators](#)

# Upsert

- Upsert is a combination of insert and **update (insert + update = upsert)**.
- We can use the upsert with different update methods.

● ● ● { upsert : true }

```
db.shoes.update(  
  {  
    brand: "Nike",  
    item: "Fragment x Sacai x Nike LDWaffle",  
    size: "UK8",  
    color: "gray"  
  },  
  { price: 18000 },  
  { upsert: true }  
)
```

# After the Class

Homework👉 *SQL vs NoSQL: comparison table*

Practice! Practice! Practice! 😊

# Welcome to the Course

- Data Modelling
- Index and Performance
- Aggregation

# Data Modeling Considerations

- The right schema design can enable you to make the most out of your database
- When designing your database it is important to consider:

Data Use and Performance

Determine important queries

Read vs Write

# Data Model Design

Embedding

vs

Referencing

user document

```
{  
  _id: <ObjectId1>,  
  username: "Chan",  
  contact: {  
    phone: "0922111444",  
    email: "chan0616@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```

user document

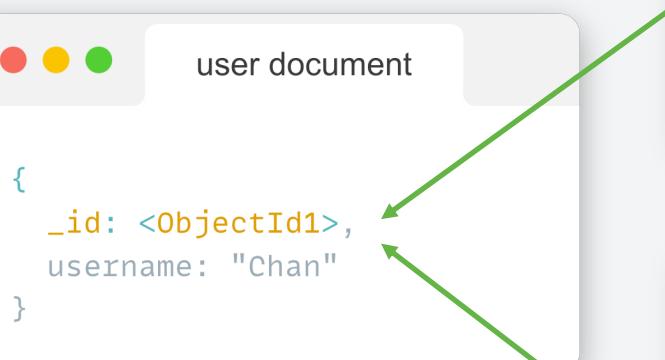
```
{  
  _id: <ObjectId1>,  
  username: "Chan"  
}
```

contact document

```
{  
  _id: <ObjectId2>,  
  user_id: <ObjectId1>,  
  phone: "0922111444",  
  email: "chan0616@example.com"  
}
```

access document

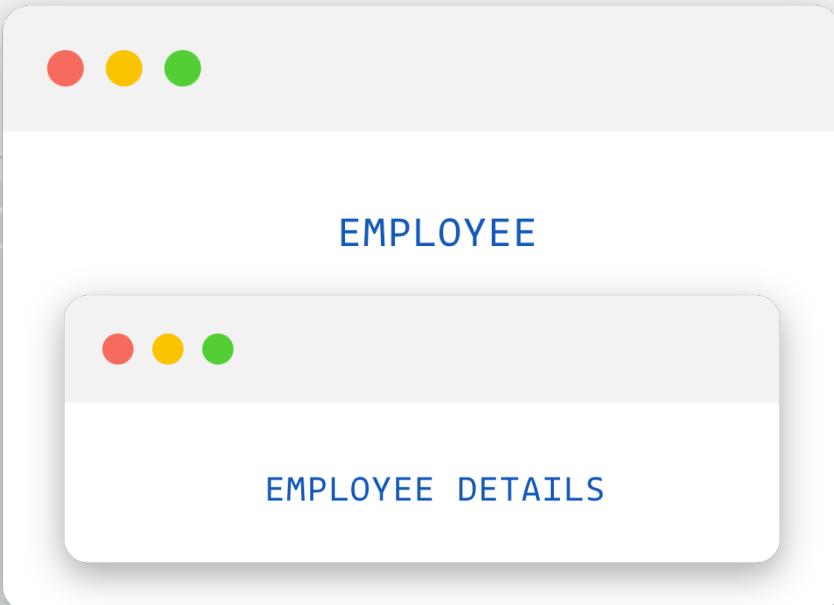
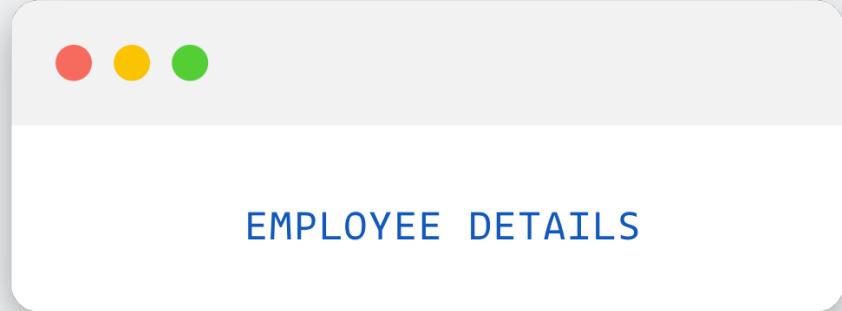
```
{  
  _id: <ObjectId3>,  
  user_id: <ObjectId1>,  
  level: 5,  
  group: "dev"  
}
```



# 1:1 Relationship

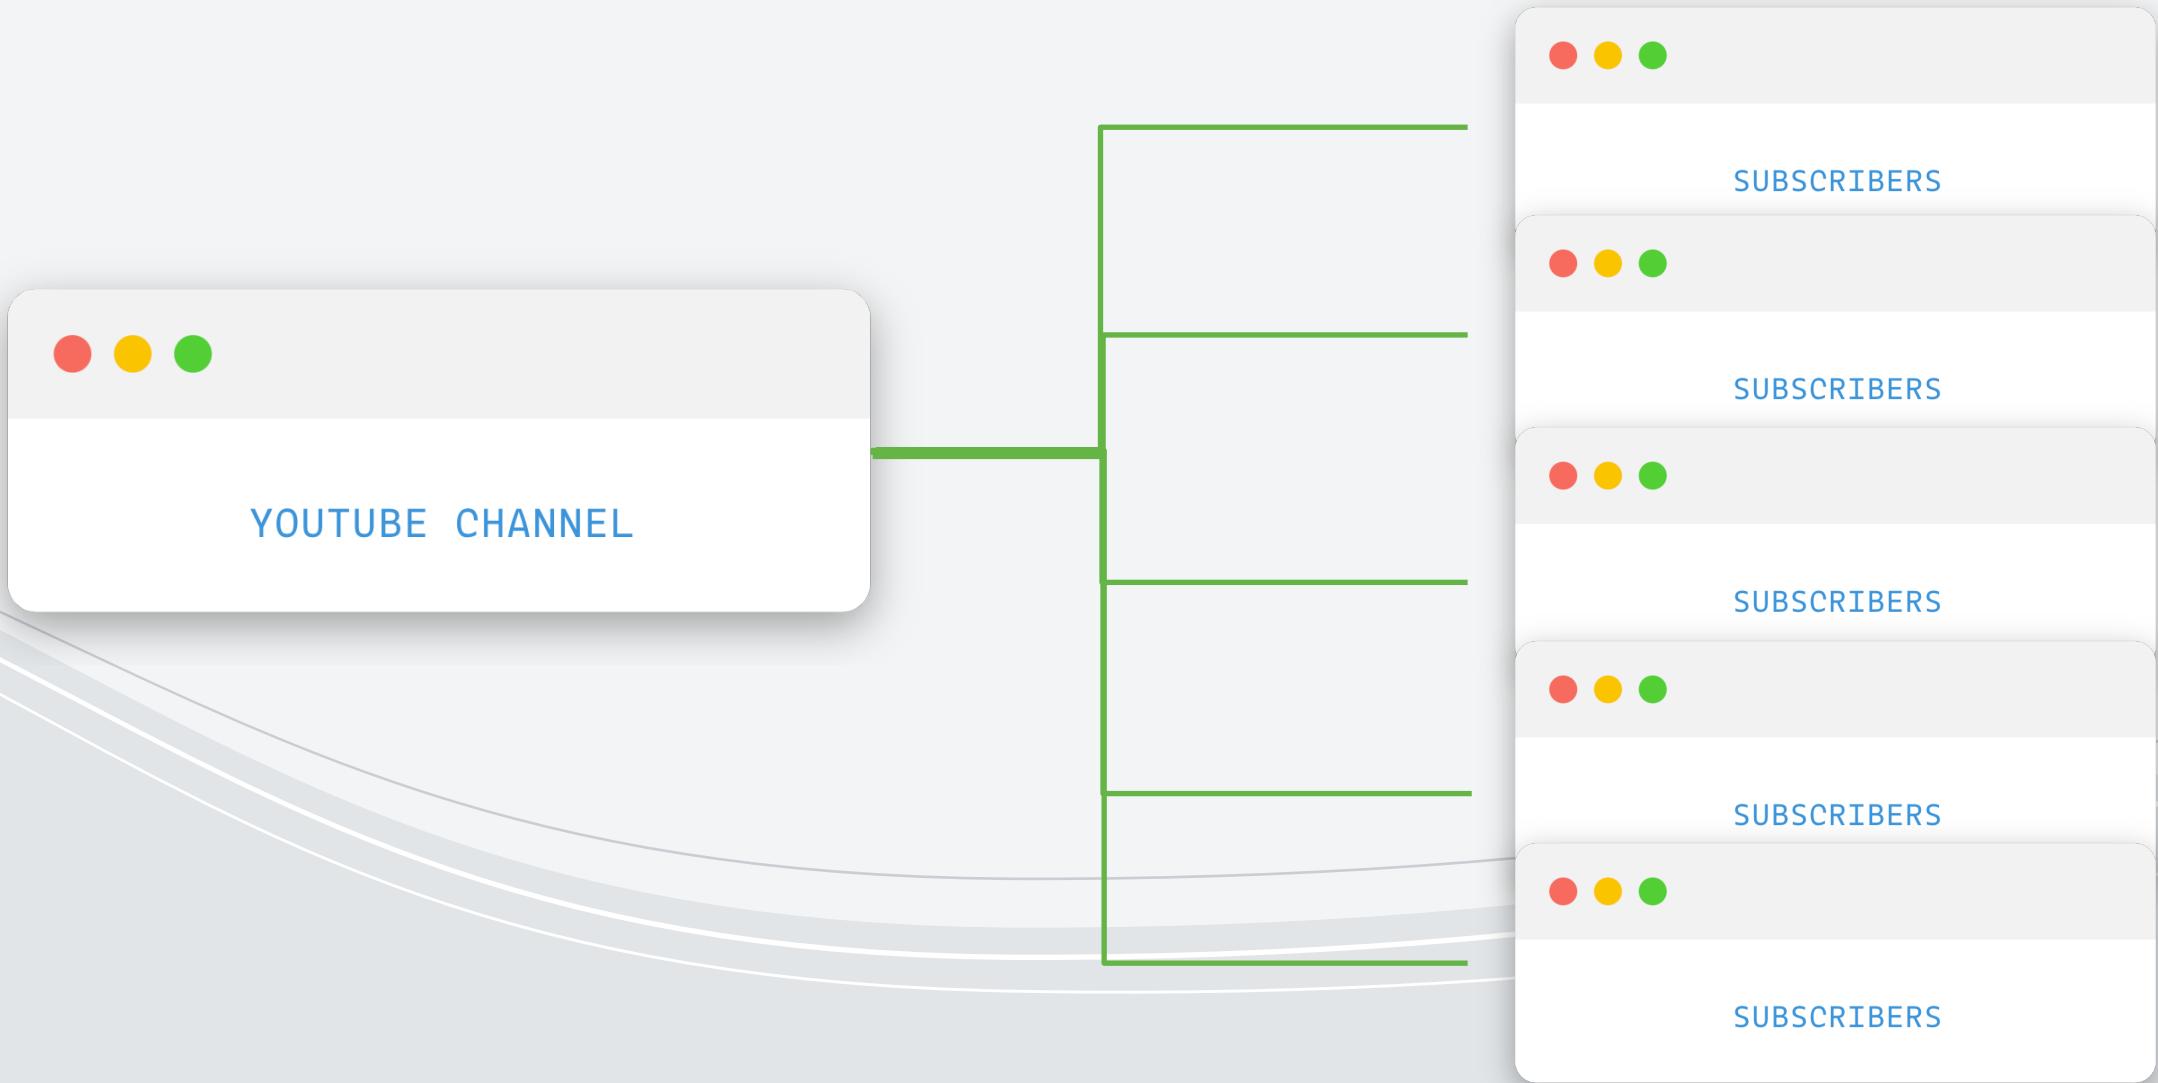


1:1

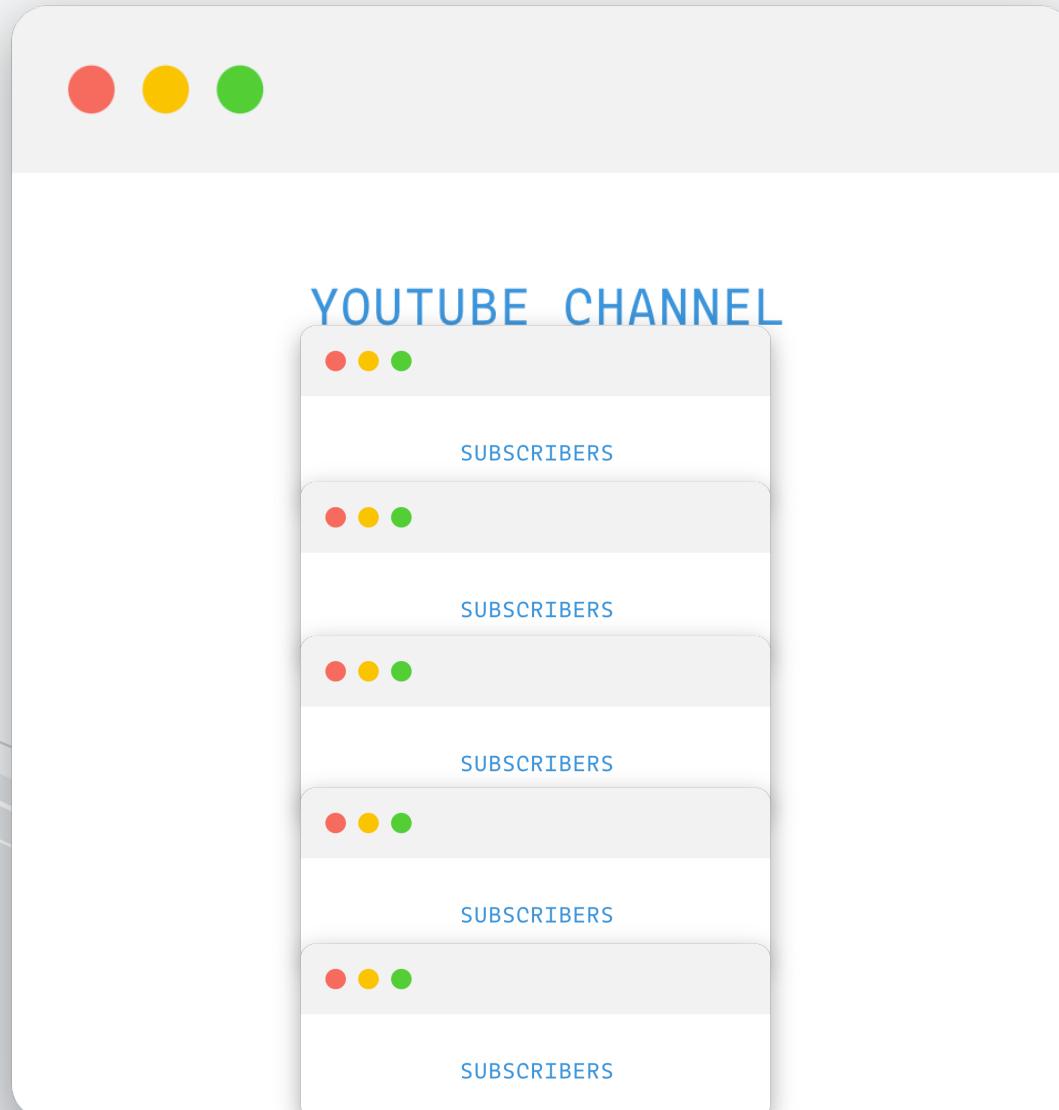


If the information of employee details is frequently retrieved then it's usually recommended that the information is embedded.

# 1:Many Relationship



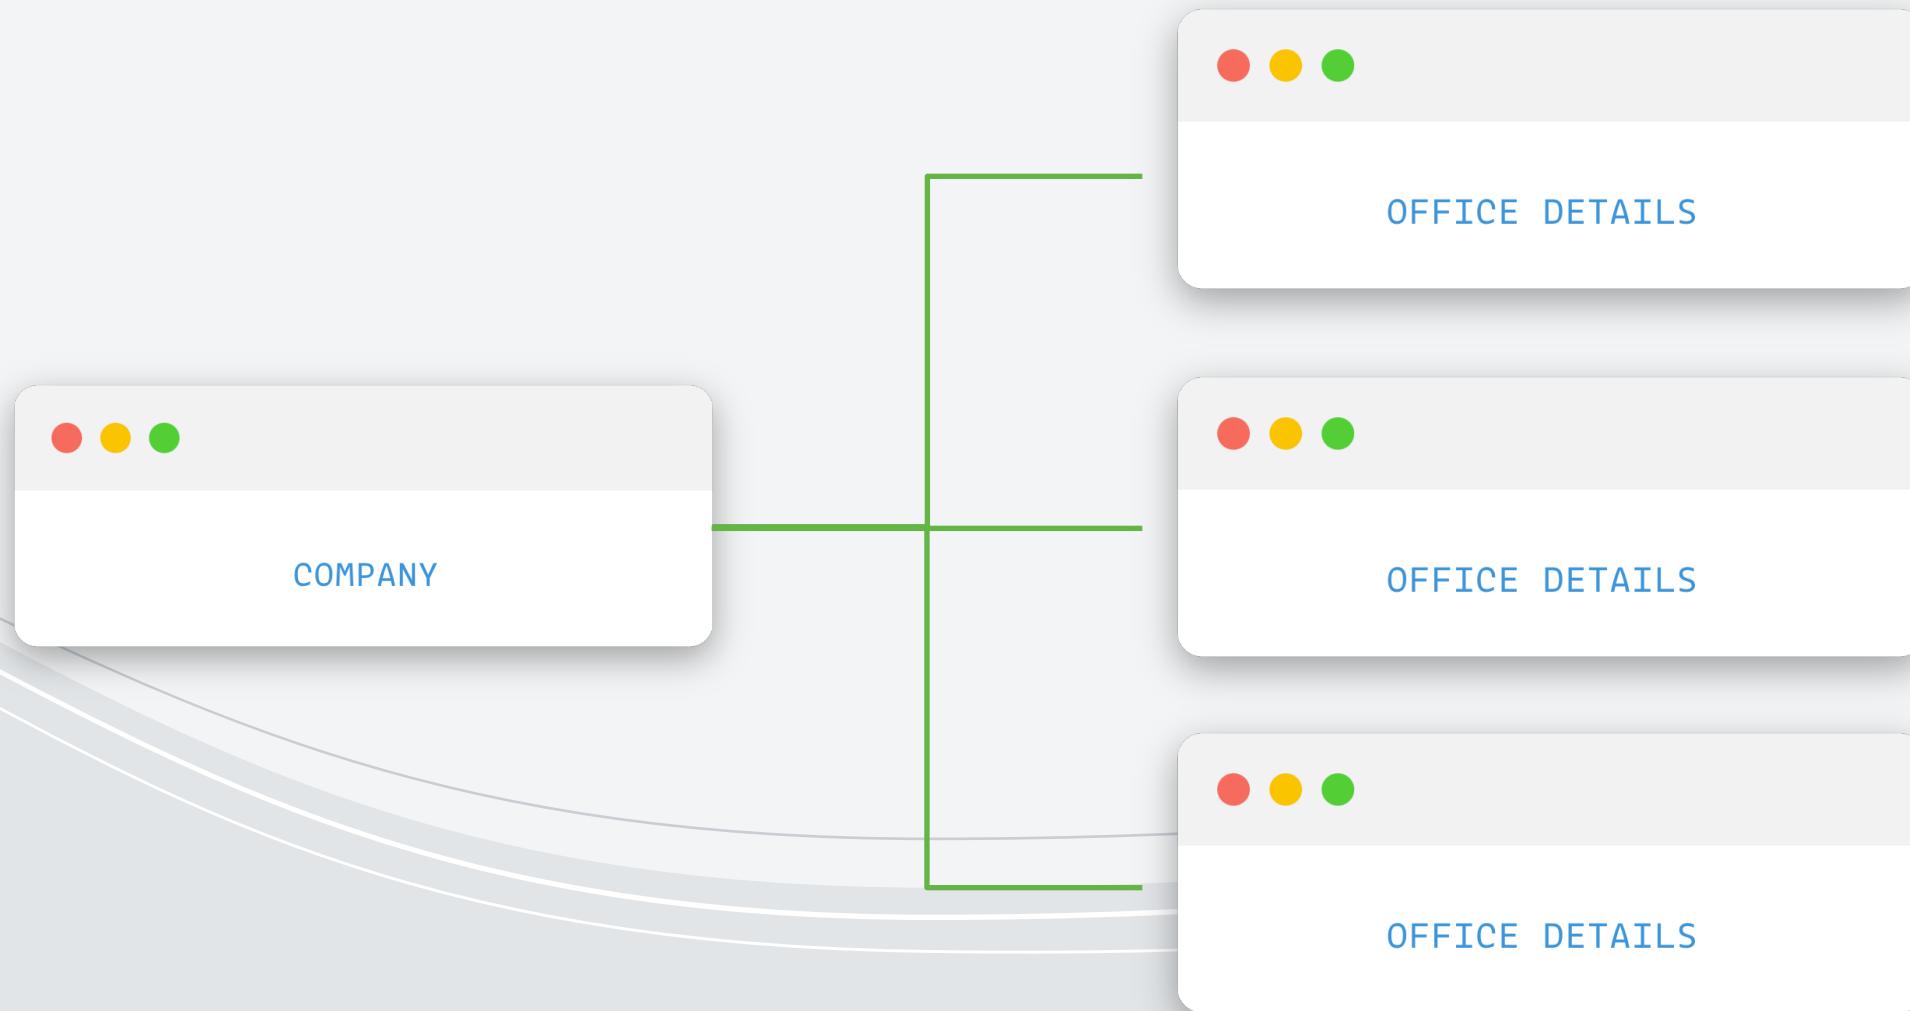
# 1:Many Relationship



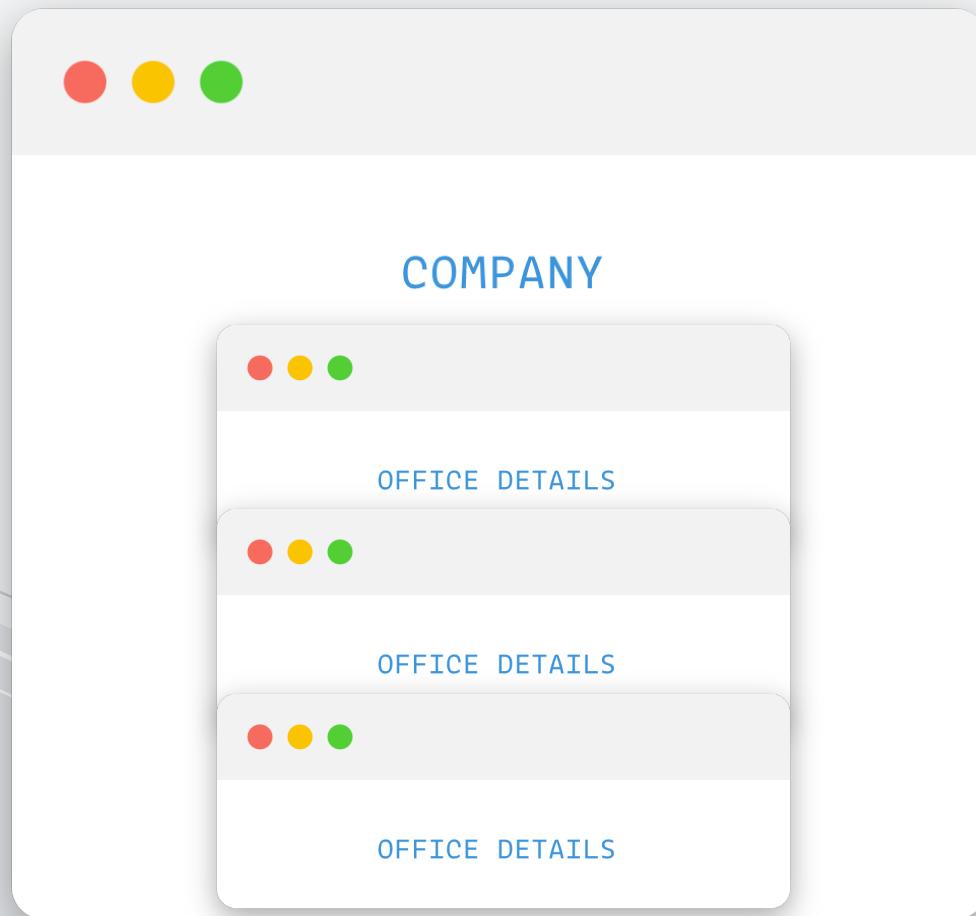
Too many embedded documents



# 1:Many Relationship

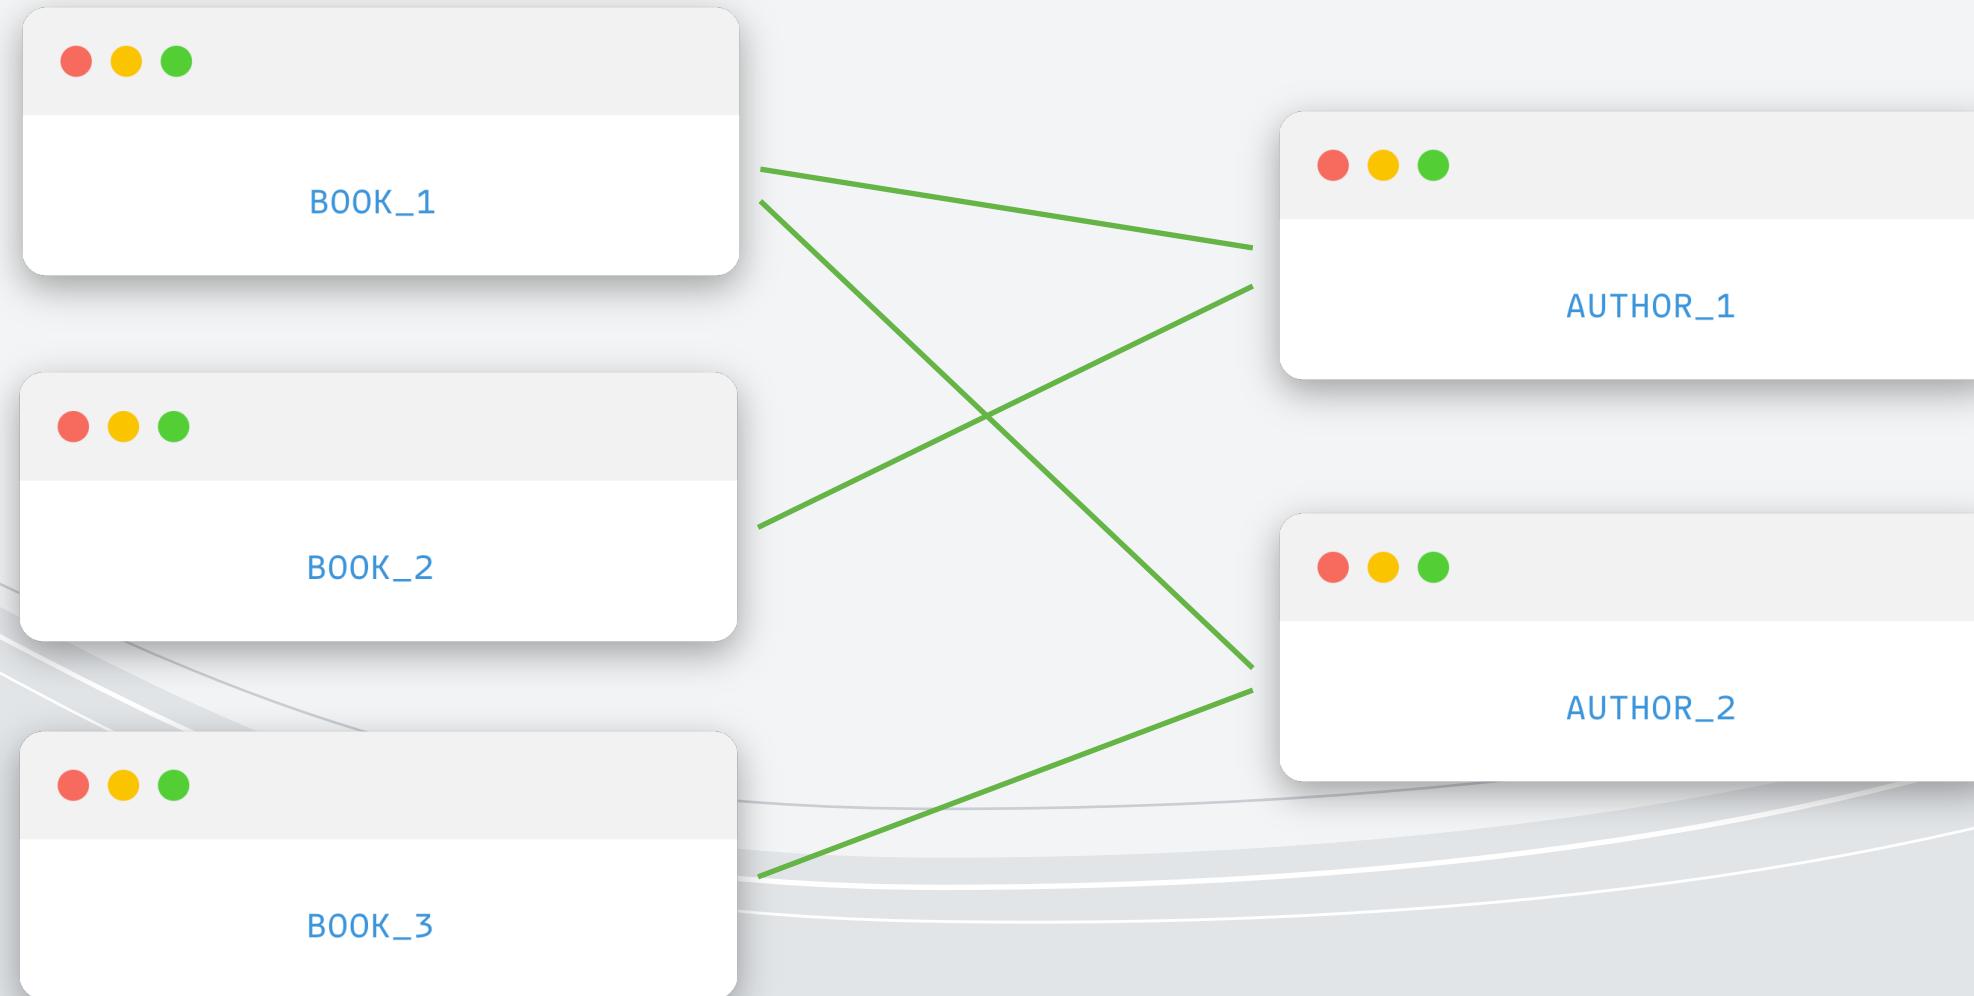


# 1:Many Relationship

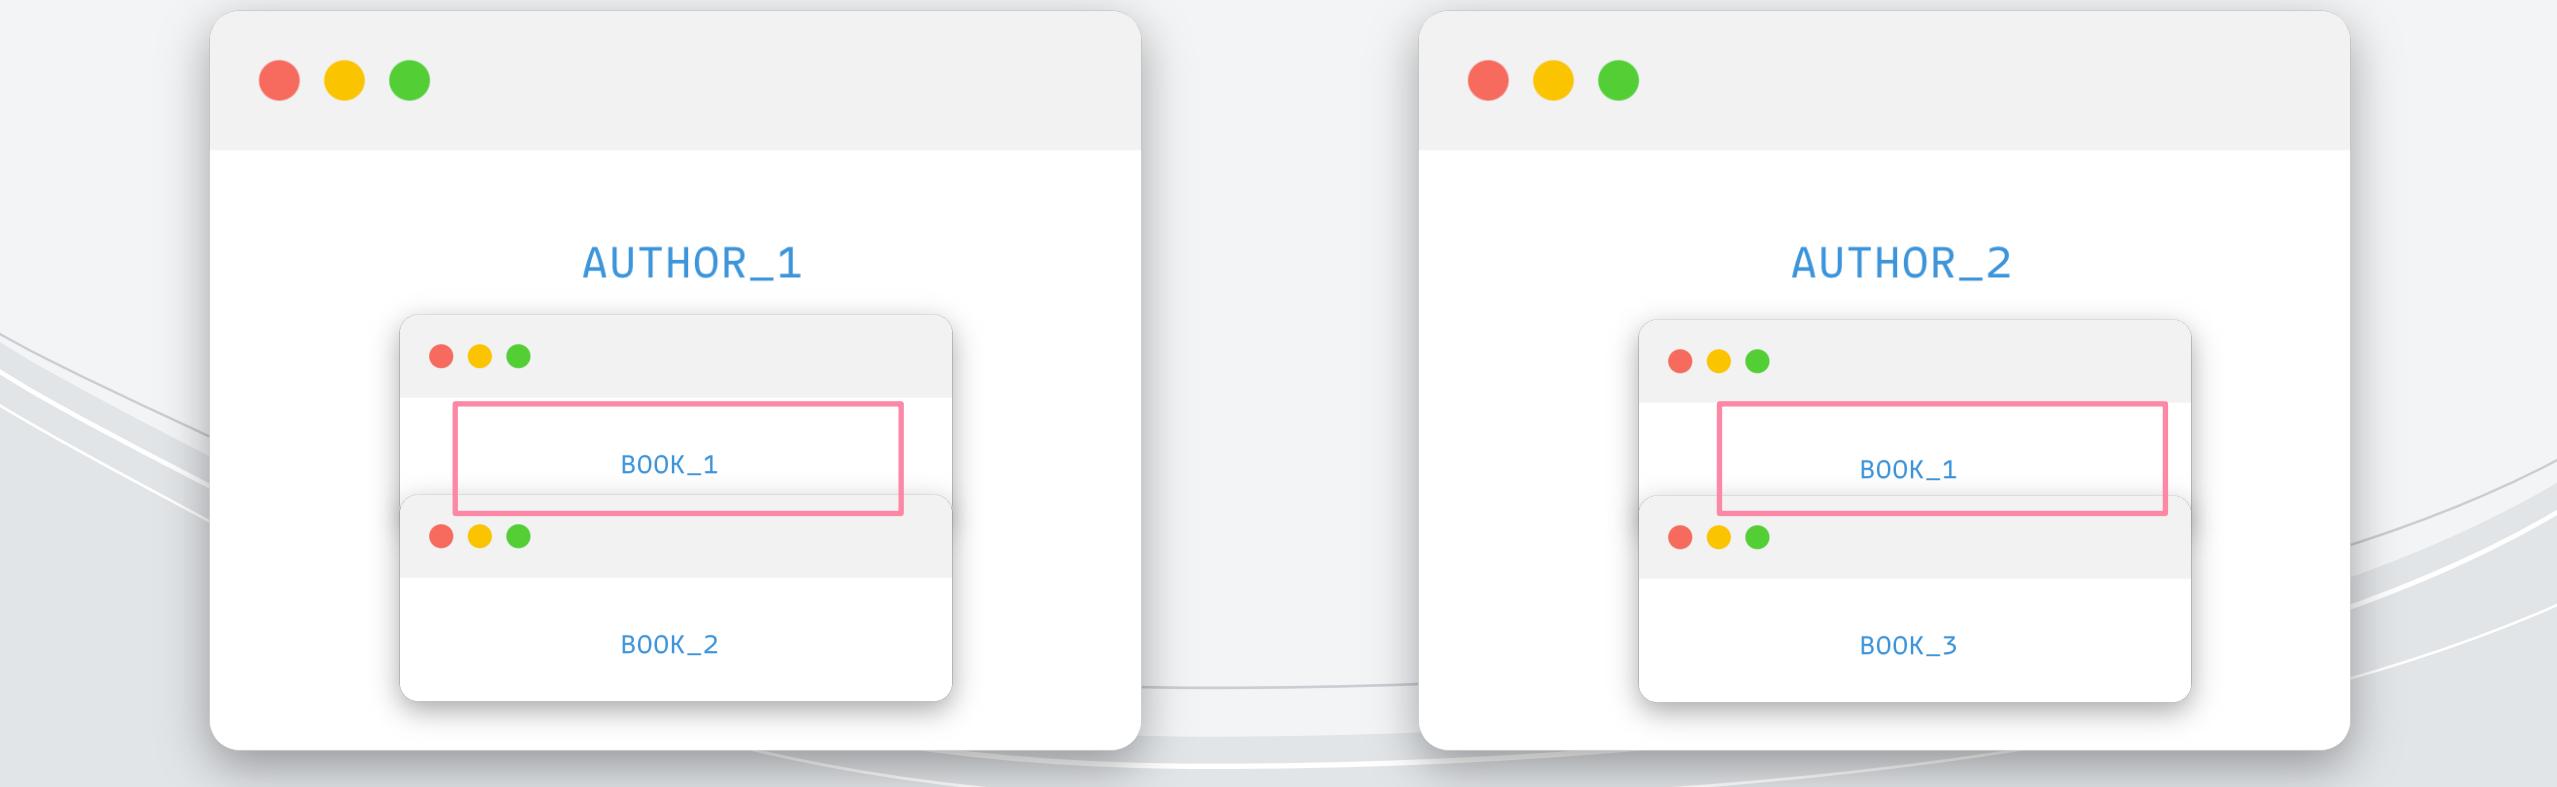


More suited to embedding documents

# Many:Many Relationship



# Many:Many Relationship



# Building with Pattern

Here is a very useful resource on [Schema Patterns](#).

Schema Patterns are very useful when deciding how to model your data in MongoDB so please be sure to check it out.

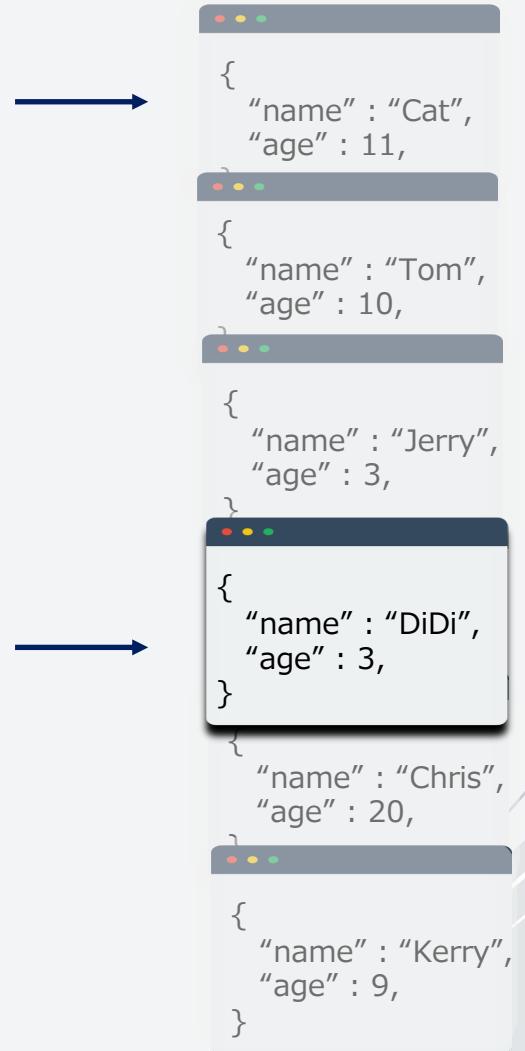
# Index and Performance

- Introduction to Index
- Single Field Index
- Compound Indexes

# Introduction to Index

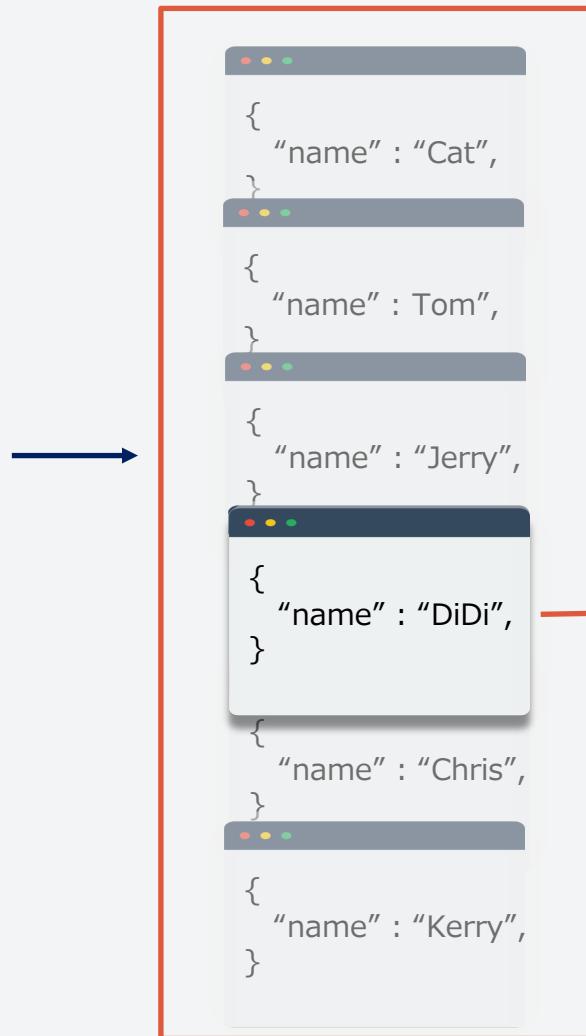
```
db.collection.find({"name":"DiDi"})
```

Collection scan(COLLSCAN) is one where every document in the collection is scanned

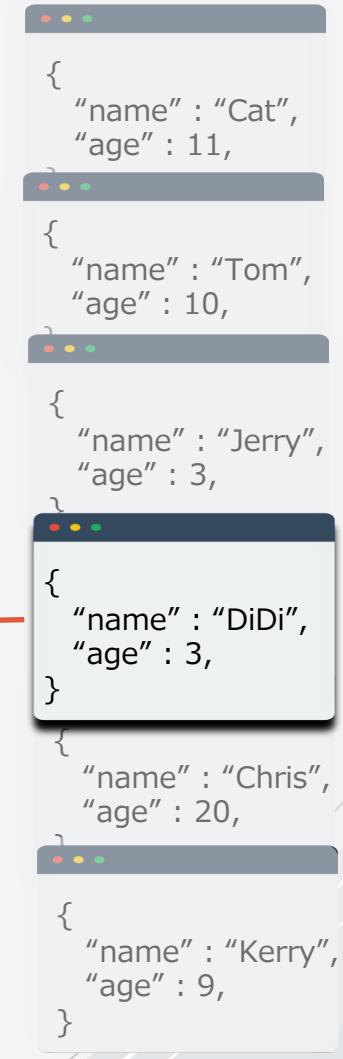


# Introduction to Index

Index on name



`db.collection.find({"name":"DiDi"})`

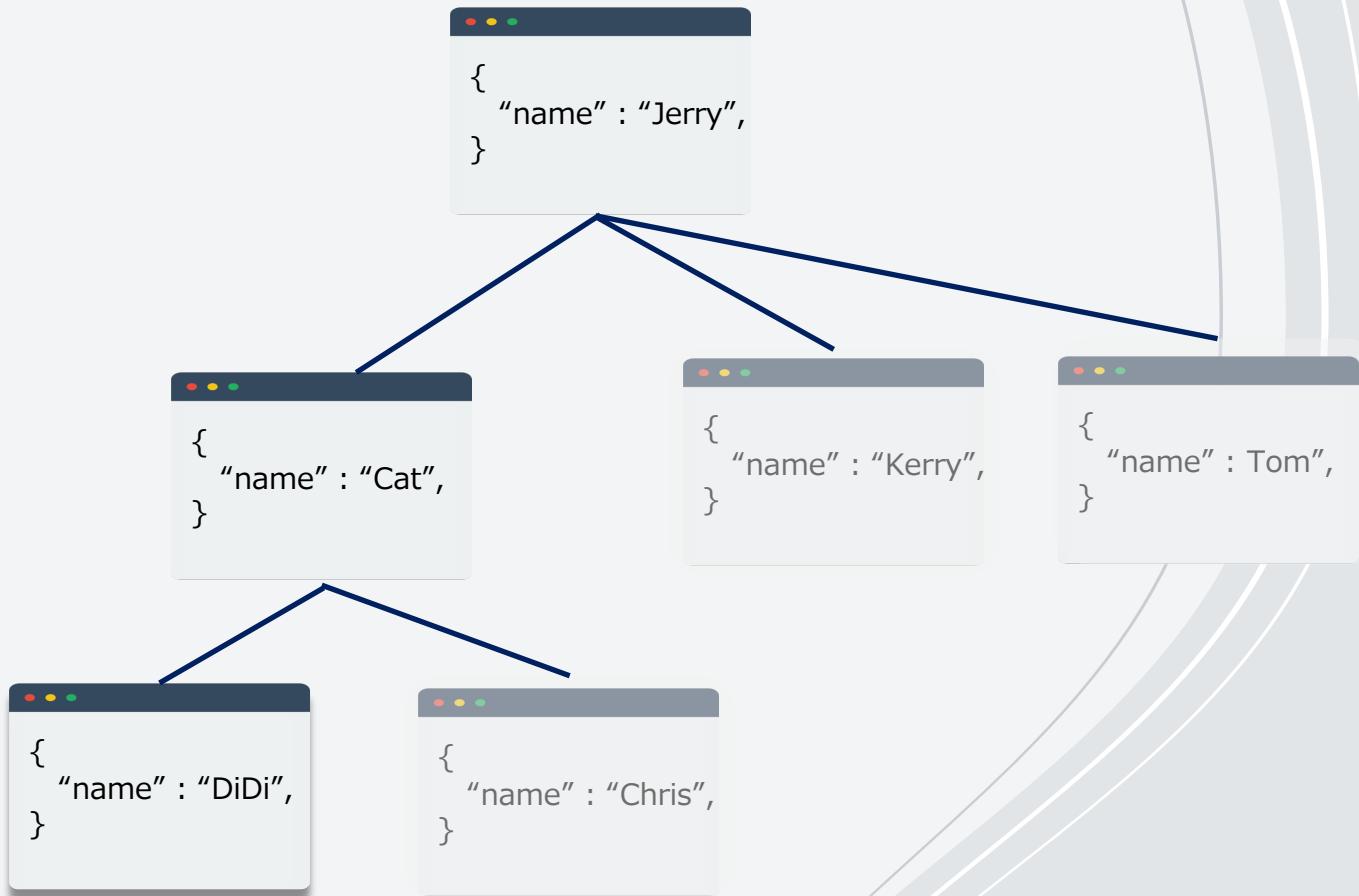


# Introduction to Index

Index on name



B-TREE



# When to use Index

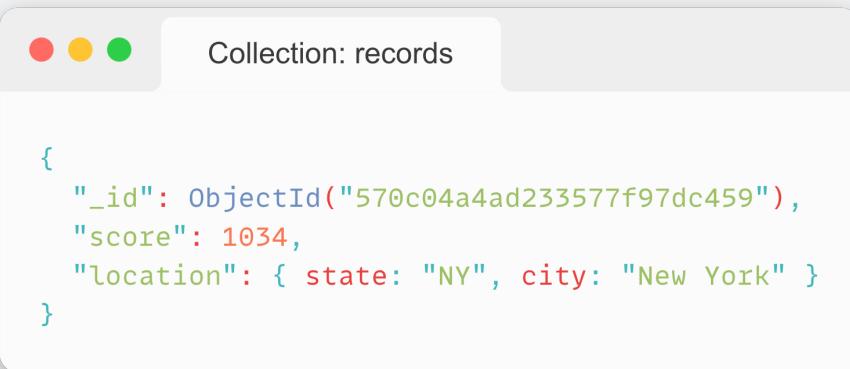
- If you have a scenario that needs sorting, when the memory sorting exceeds 32 MB.
- Certain fields have unique requirements.
- When the number of Documents is large.
- Create an Index for High-Cardinality Field.

# When is it not recommended to use Index

- It is not recommended to use Index when it cannot effectively filter data.
- Don't set Index on the field that will be updated frequently

# Single Field Index

## Create an Ascending Index on a Single Field



The MongoDB interface shows a collection named "records" with one document:

```
{  
  "_id": ObjectId("570c04a4ad233577f97dc459"),  
  "score": 1034,  
  "location": { "state": "NY", "city": "New York" }  
}
```

The following operation creates an **ascending** index on the **score** field of the **records** collection:

```
db.records.createIndex( { score: 1 } )
```

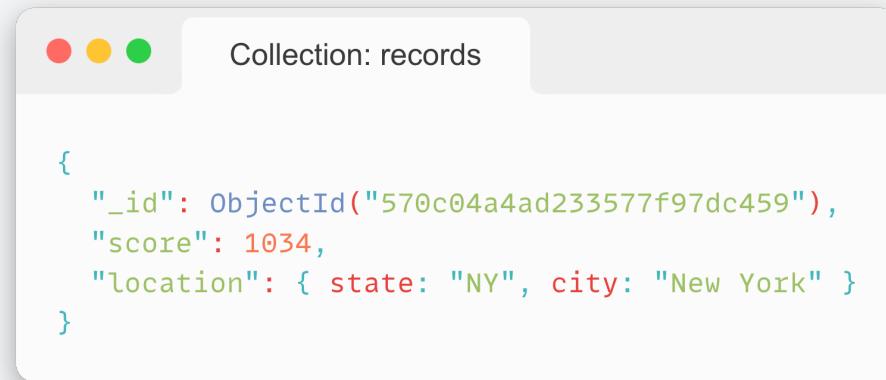
A value of **-1** specifies an index that orders items in **descending** order.

The created index will support queries that select on the field **score**, such as the following:

```
db.records.find( { score: 2 } )  
db.records.find( { score: { $gt: 10 } } )
```

# Single Field Index

## Create an Index on an Embedded Field



Collection: records

```
{  
  "_id": ObjectId("570c04a4ad233577f97dc459"),  
  "score": 1034,  
  "location": { "state": "NY", "city": "New York" }  
}
```

You can create indexes on fields within embedded documents. Indexes on embedded fields allow you to use a "dot notation," to introspect into embedded documents.

The following operation creates an index on the **location.state** field:

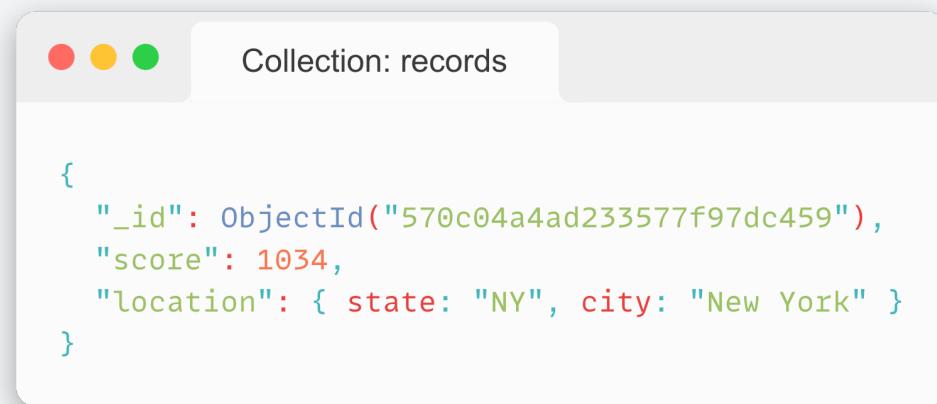
```
db.records.createIndex( { "location.state": 1 } )
```

The created index will support queries that select on the field **location.state**, such as the following:

```
db.records.find( { "location.state": "CA" } )
```

# Single Field Index

## Create an Index on Embedded Document



The screenshot shows a MongoDB interface with a collection named "records". A single document is displayed:

```
{  
  "_id": ObjectId("570c04a4ad233577f97dc459"),  
  "score": 1034,  
  "location": { "state": "NY", "city": "New York" }  
}
```

The location field is an embedded document, containing the embedded fields city and state. The following command creates an index on the location field as a whole:

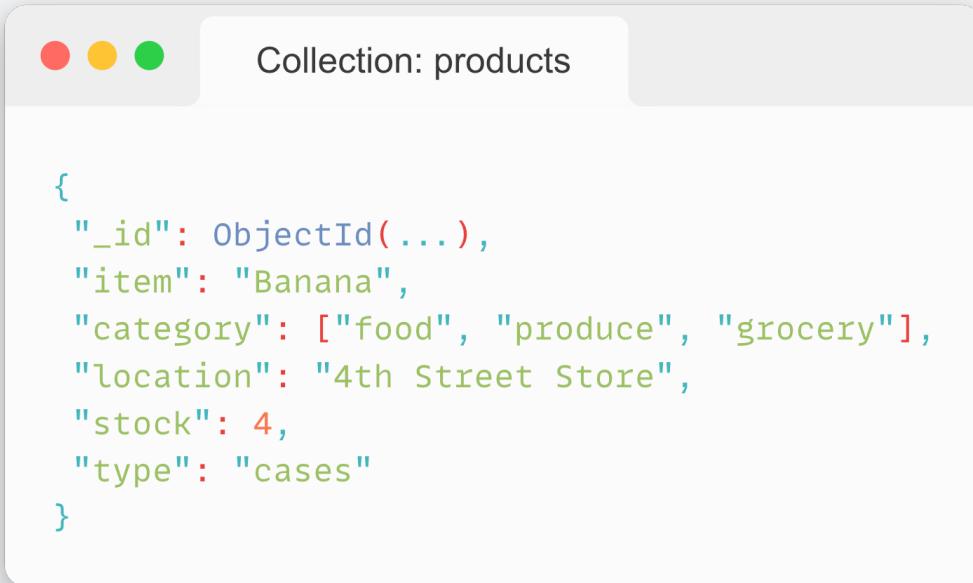
```
db.records.createIndex( { location: 1 } )
```

The following query can use the index on the **location** field:

```
db.records.find( { location: { city: "New York", state: "NY" } } )
```

# Compound Indexes

## Create a Compound Index



Collection: products

```
{  
  "_id": ObjectId(...),  
  "item": "Banana",  
  "category": ["food", "produce", "grocery"],  
  "location": "4th Street Store",  
  "stock": 4,  
  "type": "cases"  
}
```

The following operation creates an ascending index on the **item** and **stock** fields:

```
db.products.createIndex( { "item": 1, "stock": 1 } )
```

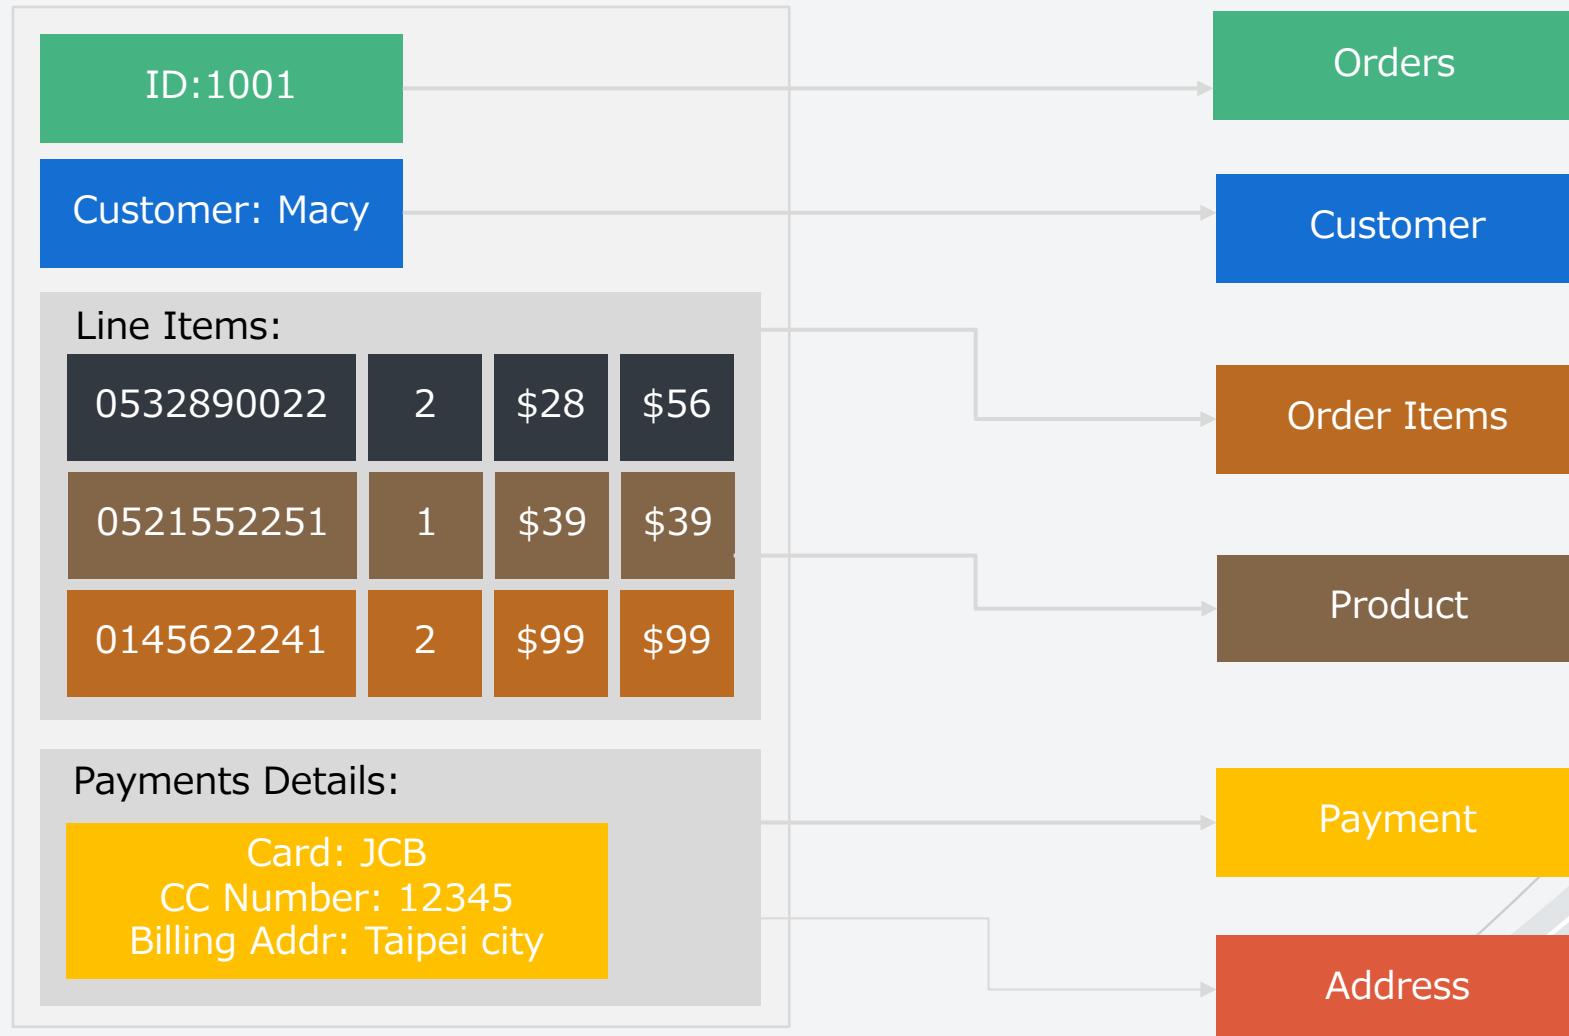
the index supports queries on the item field as well as both item and stock fields:

```
db.products.find( { item: "Banana" } )  
db.products.find( { item: "Banana", stock: { $gt: 5 } } )
```

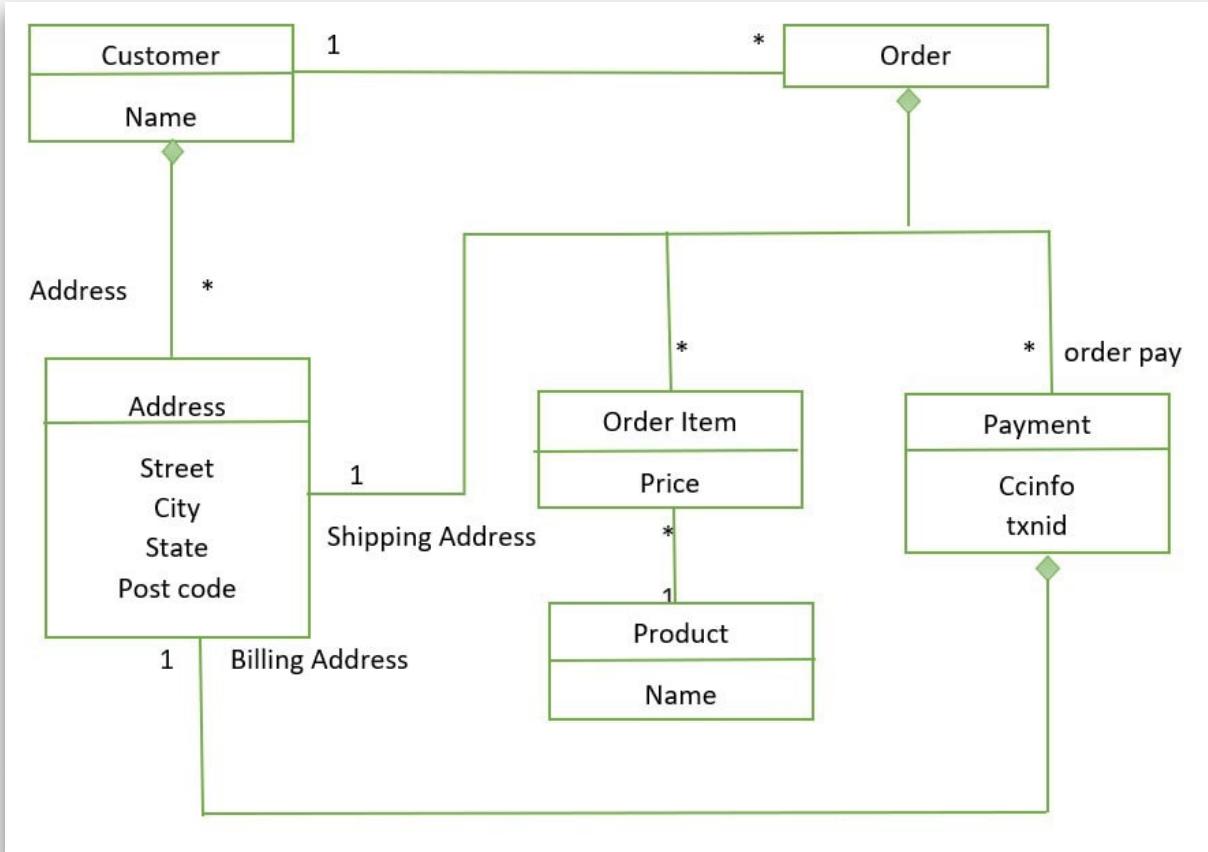
# How to design and use Index

- Use Compound Index instead of Single Field Index
- Follow ESR rule
  - Equality : Field for equality query.
  - Sort : Field to sort.
  - Range : Field of the data range to access.

# Aggregate Data Models



# Aggregate Data Models



Computed Results

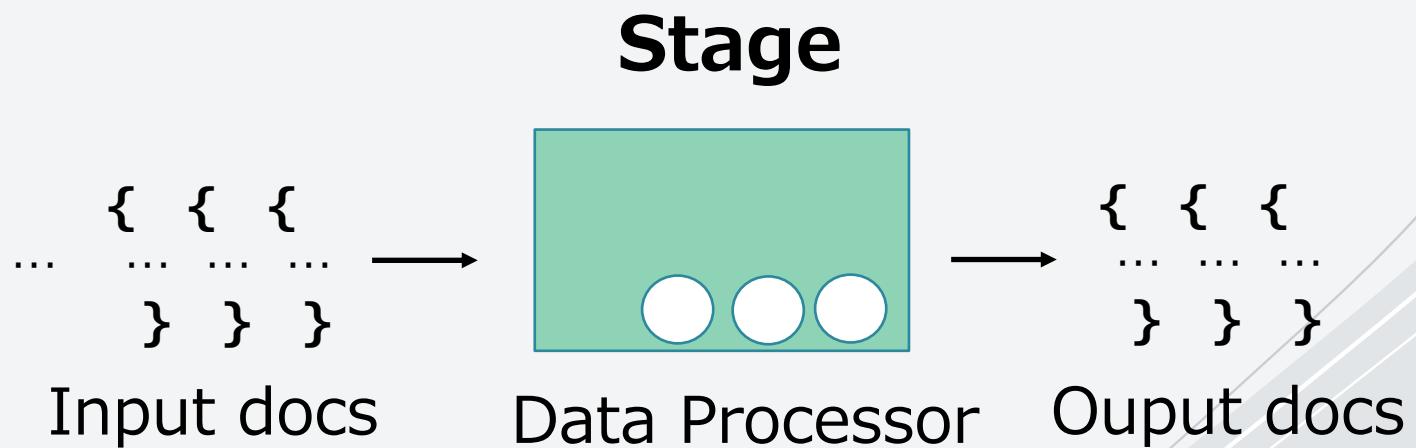
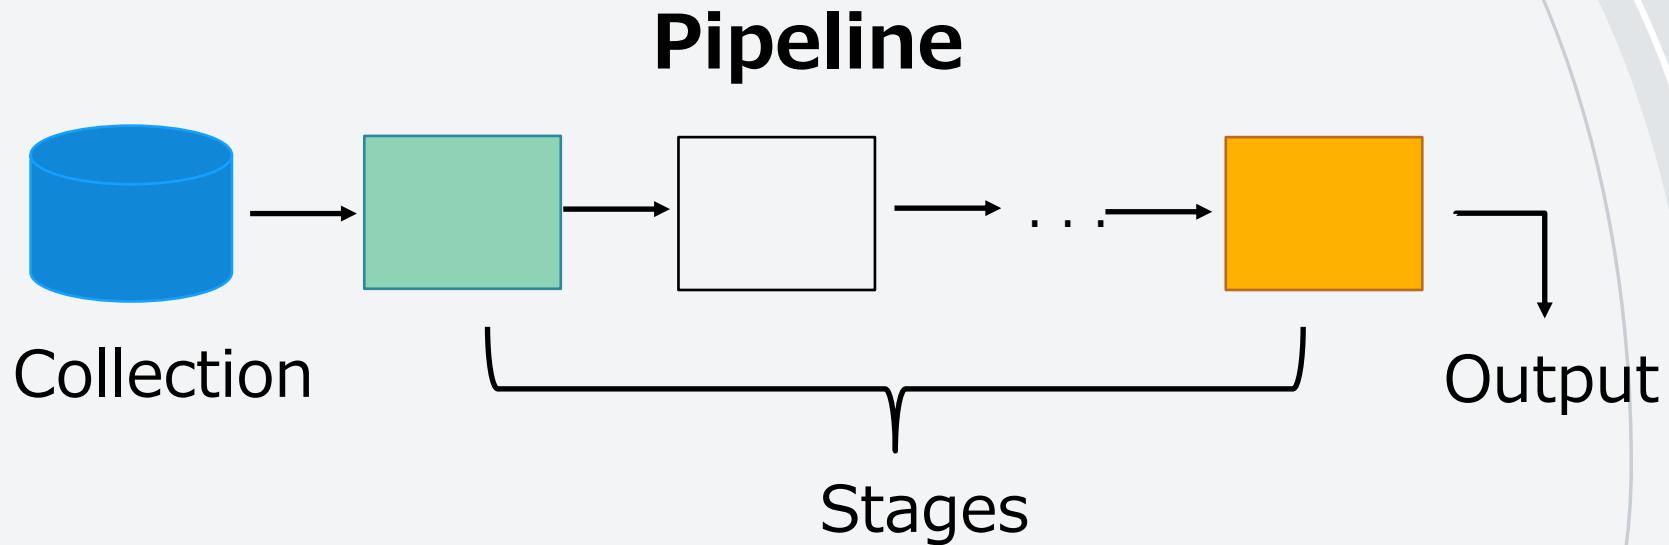
```
{
  "customer": {
    "id": 1,
    "name": "Martin",
    "billingAddress": [
      {
        "city": "Chicago"
      }
    ],
    "orders": [
      {
        "id": 99,
        "customerId": 1,
        "orderItems": [
          {
            "productId": 27,
            "price": 32.45,
            "productName": "NoSQL Distilled"
          }
        ],
        "shippingAddress": [
          {
            "city": "Chicago"
          }
        ]
      }
    ],
    "orderPayment": [
      {
        "ccinfo": "1000-1000-1000-1000",
        "txnid": "abelif879rft",
        "billingAddress": {
          "city": "Chicago"
        }
      }
    ],
    "payments": []
  }
}
```

# Aggregation Pipeline Example

Aggregation operations process data records and return computed results.

```
aggregate( [  
    // Stage 1: Filter pizza order documents by pizza size  
    {  
        $match: { size: "medium" }  
    },  
  
    // Stage 2: Group remaining documents by pizza name and calculate total quantity  
    {  
        $group: { _id: "$name", totalQuantity: { $sum: "$quantity" } }  
    }  
]
```

# Aggregation Pipelines



# Aggregation Pipeline Stages/Operators

## 👉 Aggregation Pipeline Stages

- **\$lookup**
- \$match
- \$group
- \$set
- \$unset
- \$gt

## 👉 Aggregation Pipeline Operators

# \$lookup

## Equality Match with a Single Join Condition

\$lookup

```
{  
  $lookup:  
    {  
      from: <collection to join>,  
      localField: <field from the input documents>,  
      foreignField: <field from the documents of the "from" collection>,  
      as: <output array field>  
    }  
}
```

pseudo-SQL

```
SELECT *, <output array field>  
FROM collection  
WHERE <output array field> IN (  
  SELECT *  
  FROM <collection to join>  
  WHERE <foreignField> = <collection.localField>  
) ;
```

# Perform a Single Equality Join with \$lookup



Collection: orders

```
db.orders.insertMany( [  
    { "_id" : 1, "item" : "almonds", "price" : 12, "quantity" : 2 },  
    { "_id" : 2, "item" : "pecans", "price" : 20, "quantity" : 1 },  
    { "_id" : 3 }  
] )
```



Collection: inventory

```
db.inventory.insertMany( [  
    { "_id" : 1, "sku" : "almonds", "description": "product 1", "instock" : 120 },  
    { "_id" : 2, "sku" : "bread", "description": "product 2", "instock" : 80 },  
    { "_id" : 3, "sku" : "cashews", "description": "product 3", "instock" : 60 },  
    { "_id" : 4, "sku" : "pecans", "description": "product 4", "instock" : 70 },  
    { "_id" : 5, "sku": null, "description": "Incomplete" },  
    { "_id" : 6 }  
] )
```



\$lookup

```
db.orders.aggregate( [  
    {  
        $lookup:  
            {  
                from: "inventory",  
                localField: "item",  
                foreignField: "sku",  
                as: "inventory_docs"  
            }  
    }  
] )
```



pseudo-SQL

```
SELECT *, inventory_docs  
FROM orders  
WHERE inventory_docs IN (  
    SELECT *  
    FROM inventory  
    WHERE sku = orders.item  
);
```

# \$lookup

## Correlated Subqueries Using Concise Syntax

```
$lookup  
{  
  $lookup:  
    {  
      from: <foreign collection>,  
      localField: <field from local collection's documents>,  
      foreignField: <field from foreign collection's documents>,  
      let: { <var_1>: <expression>, ..., <var_n>: <expression> },  
      pipeline: [ <pipeline to run> ],  
      as: <output array field>  
    }  
}  
}
```

```
pseudo-SQL  
  
SELECT *, <output array field>  
FROM localCollection  
WHERE <output array field> IN (  
  SELECT <documents as determined from the pipeline>  
  FROM <foreignCollection>  
  WHERE <foreignCollection.foreignField> = <localCollection.localField>  
  AND <pipeline match condition>  
);
```

👉 The \$lookup accepts a document with fields

# \$lookup example: pipeline example with condition

**let (optional)**: To reference variables in pipeline stages, use the "`$$<variable>`" syntax.

**pipeline**: Specifies the pipeline to run on the joined collection. The pipeline determines the resulting documents from the joined collection.

**\$expr**: \$expr can build query expressions that compare fields from the same document in a \$match stage. If the \$match stage is part of a \$lookup stage, \$expr can compare fields using let variables.

```
Aggregation Pipeline

db.orders.aggregate( [
  {
    $lookup: {
      from: "restaurants",
      let: { orders_restaurant_name: "$restaurant_name",
             orders_drink: "$drink" },
      pipeline: [ {
        $match: {
          $expr: {
            $and: [
              { $eq: [ "$$orders_restaurant_name", "$name" ] },
              { $in: [ "$$orders_drink", "$beverages" ] }
            ]
          }
        }
      } ],
      as: "matches"
    }
  }
])
```

# Perform a Concise Correlated Subquery with \$lookup

```
Collection: restaurants
  db.orders.insertMany([
    db.restaurants.insertMany([
      {
        _id: 1,
        name: "American Steak House",
        food: [ "filet", "sirloin" ],
        beverages: [ "beer", "wine" ],
        item: "cheese pizza",
        restaurant_name: "Honest John Pizza",
        drink: "lemonade"
      },
      {
        _id: 2,
        name: "Honest John Pizza",
        food: [ "cheese pizza", "pepperoni pizza" ],
        beverages: [ "soda" ],
        item: "cheese pizza",
        restaurant_name: "Honest John Pizza",
        drink: "soda"
      }
    ]
  ])
}

Collection: orders
  db.orders.insertMany([
    {
      _id: 1,
      restaurant_name: "American Steak House",
      drink: "beer"
    },
    {
      _id: 2,
      restaurant_name: "Honest John Pizza",
      drink: "lemonade"
    },
    {
      _id: 3,
      restaurant_name: "Honest John Pizza",
      drink: "soda"
    }
  ])

```

Aggregation Pipeline

```
db.orders.aggregate( [
  {
    $lookup: {
      from: "restaurants",
      let: { orders_restaurant_name: "$restaurant_name",
             orders_drink: "$drink" },
      pipeline: [ {
        $match: {
          $expr: {
            $and: [
              { $eq: [ "$$orders_restaurant_name", "$name" ] },
              { $in: [ "$$orders_drink", "$beverages" ] }
            ]
          }
        }
      } ],
      as: "matches"
    }
  }
] )
```

pseudo-SQL

```
SELECT *, matches
FROM orders
WHERE matches IN (
  SELECT *
  FROM restaurants
  WHERE restaurants.name = orders.restaurant_name
  AND restaurants.beverages = orders.drink
);
```

# Aggregation - Multi-Field Join & One-to-Many

## Scenario

You want to generate a report to list all the orders made for each product in 2020. To achieve this, you need to take a shop's **products collection** and join each product record to all its orders stored in an **orders collection**. There is a 1:many relationship between both collections, based on a match of two fields on each side. Rather than joining on a single field like **product\_id** (which doesn't exist in this data set), you need to use two common fields to join (**product\_name** and **product\_variation**). Note that the requirement to **perform a 1:many join** does not mandate the need to join by multiple fields on each side of the join. However, in this example, it has been deemed beneficial to show both of these aspects in one place.



[Exercises: 01\\_aggregation-oneToMany](#)

# Sample Data

Collection: products

```
db.products.insertMany([
  {
    "name": "Asus Laptop",
    "variation": "Ultra HD",
    "category": "ELECTRONICS",
    "description": "Great for watching movies",
  },
  {
    "name": "Asus Laptop",
    "variation": "Normal Display",
    "category": "ELECTRONICS",
    "description": "Good value laptop for students",
  },
  {
    "name": "The Day Of The Triffids",
    "variation": "1st Edition",
    "category": "BOOKS",
    "description": "Classic post-apocalyptic novel",
  },
  {
    "name": "The Day Of The Triffids",
    "variation": "2nd Edition",
    "category": "BOOKS",
    "description": "Classic post-apocalyptic novel",
  },
  {
    "name": "Morphy Richards Food Mixer",
    "variation": "Deluxe",
    "category": "KITCHENWARE",
    "description": "Luxury mixer turning good cakes into great",
  },
  {
    "name": "Karcher Hose Set",
    "variation": "Full Monty",
    "category": "GARDEN",
    "description": "Hose + nosels + winder for tidy storage",
  },
]);
])
```

Collection: orders

```
// Create index for the orders collection
db.orders.createIndex({ "product_name": 1, "product_variation": 1 });

// Insert 4 records into the orders collection
db.orders.insertMany([
  {
    "customer_id": "elise_smith@myemail.com",
    "orderdate": ISODate("2020-05-30T08:35:52Z"),
    "product_name": "Asus Laptop",
    "product_variation": "Normal Display",
    "value": NumberDecimal("431.43"),
  },
  {
    "customer_id": "tj@wheresmyemail.com",
    "orderdate": ISODate("2019-05-28T19:13:32Z"),
    "product_name": "The Day Of The Triffids",
    "product_variation": "2nd Edition",
    "value": NumberDecimal("5.01"),
  },
  {
    "customer_id": "oranieri@wammmail.com",
    "orderdate": ISODate("2020-01-01T08:25:37Z"),
    "product_name": "Morphy Richards Food Mixer",
    "product_variation": "Deluxe",
    "value": NumberDecimal("63.13"),
  },
  {
    "customer_id": "jjones@tepidmail.com",
    "orderdate": ISODate("2020-12-26T08:55:46Z"),
    "product_name": "Asus Laptop",
    "product_variation": "Normal Display",
    "value": NumberDecimal("429.65"),
  },
]);
])
```

```
var pipeline = [
  // Join by 2 fields in products collection to 2 fields
  // in orders collection
  {"$lookup": {
    "from": "orders",
    "let": {
      "prdname": "$name",
      "prdvartn": "$variation",
    },
    // Embedded pipeline to control how the join is matched
    "pipeline": [
      // Join by two fields in each side
      {"$match": {
        "$expr": {
          "$and": [
            {"$eq": ["$product_name", "$$prdname"]},
            {"$eq": ["$product_variation", "$$prdvartn"]}
          ]
        }
      }},
      // Match only orders made in 2020
      {"$match": {
        "orderdate": {
          "$gte": ISODate("2020-01-01T00:00:00Z"),
          "$lt": ISODate("2021-01-01T00:00:00Z")
        }
      }},
      // Exclude some unwanted fields from the right side of the join
      {"$unset": [
        "_id",
        "product_name",
        "product_variation",
      ]},
      ],
      as: "orders",
    },
    // Only show products that have at least one order
    {"$match": {
      "orders": {"$ne": []},
    }},
    // Omit unwanted fields
    {"$unset": [
      "_id",
    ]}
]
```

# Execution

```
db.products.aggregate(pipeline);
```



## MongoDB Exercises